

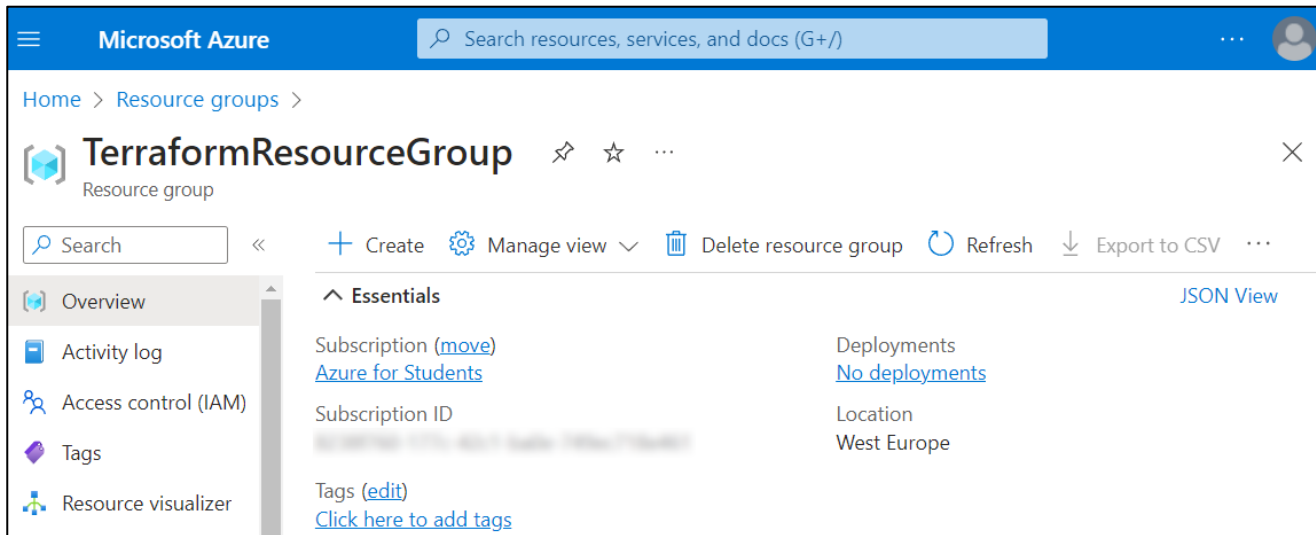
Exercise: IaC and Monitoring

Exercise assignment for the ["Containers and Clouds" course @ SoftUni](#).

I. Infrastructure as Code

1. Azure Resource Group

Now you have a task to **create a Terraform configuration to deploy an Azure resource group**.



Hints

Open a **terminal** (for example PowerShell), **create a Terraform configuration folder** with an **empty configuration file** and **follow the steps below** to fulfill the task:

1. **Authenticate** using the **Azure CLI**, i.e. **log in to Azure**, as **Terraform must authenticate** to create infrastructure
2. **Write the configuration** for creating an **Azure resource group**
 - You need an **Azure provider**, available here:
<https://registry.terraform.io/providers/hashicorp/azurerm/latest>
 - The **Azure provider** needs a **feature { }** block in the **configuration**
 - At the end, the **resource group** should be created using the **"azurerm_resource_group"** **Terraform resource**, whose **required arguments** can be seen here:
https://registry.terraform.io/providers/hashicorp/azurerm/latest/docs/resources/resource_group

The **configuration file** looks like shown below. The **resource group name and location** are for you to choose:

```
azure-rg.tf
D: > SoftUni > azure-resource-group > azure-rg.tf
1 terraform {
2   provider "azurerm" {
3     subscription_id = "xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx"
4     client_id       = "xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx"
5     client_secret    = "xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx"
6   }
7 }
8 }
9
10 provider "azurerm" {
11   features {}
12 }
13
14 resource "azurerm_resource_group" "rg" {
15   name     = "ContactBook-rg"
16   location = "West Europe"
17 }
```

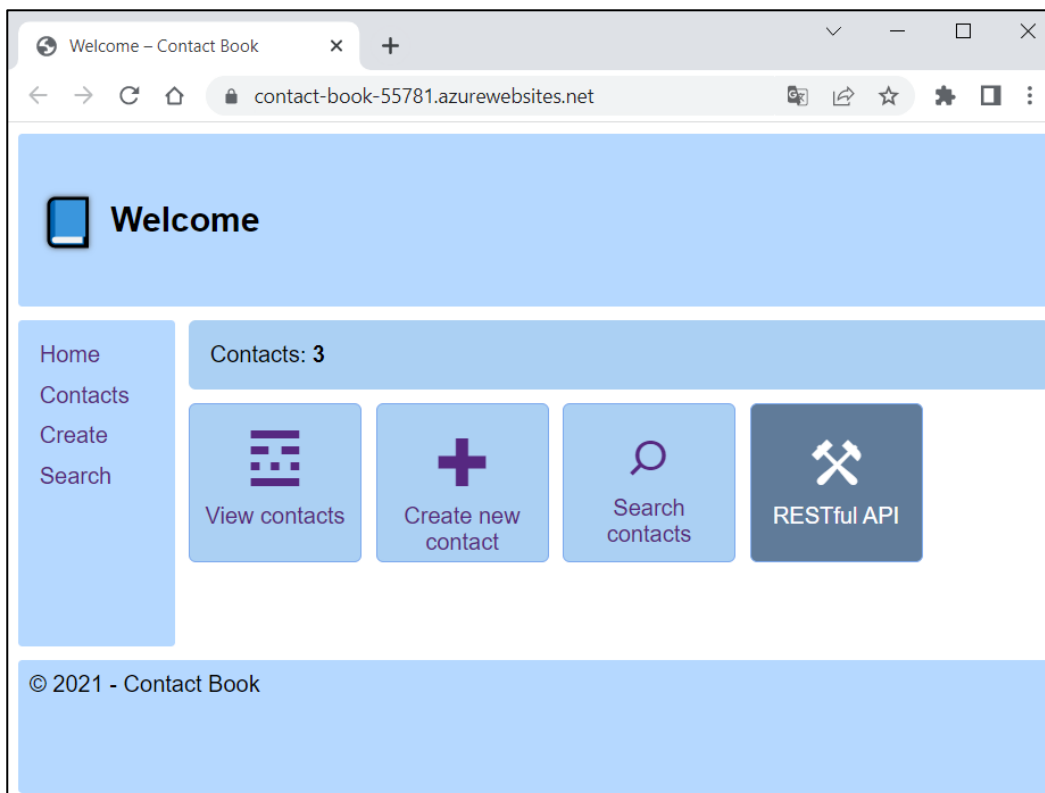
3. Initialize, format, validate and apply your Terraform configuration
4. Navigate to **Azure Portal** in the browser and validate that a **resource group** was created

Later you can **delete the resource group** from **Azure** again using **Terraform**.

As we know how to **create an Azure resource group with Terraform**, let's see how this would be **useful for us in the next task**.

2. Azure Web App

You are already **familiar with Azure Web Apps** and now you should **use Terraform** to **create a resource group**, then **create an App Service Plan** and finally **deploy the "Contact Book" app** to **Azure** from a **GitHub repo**.



"Contact Book" is a **Node.js app without a database**, available here: <https://github.com/nakov/ContactBook>.

Hints

To fulfill your task, you need to create a Terraform configuration file. Find the Terraform resources you need in the Terraform Registry and use them: <https://registry.terraform.io>.

The configuration you should write:

- Uses and configures an **Azure provider** (as in the previous exercise)

```
azure-app.tf ●
D: > SoftUni > azure-app-deploy > azure-app.tf
1  # Configure the Azure provider
2
3
4
5
6
7
8
9
10
11
12
13
14
```

- Generates a random integer with minimum and maximum number range to be used for creating unique resource names

```
15  # Generate a random integer to create a globally unique name
16  resource "random_integer" "ri" {
17      min = 10000
18      max = 99999
19  }
20
```

- Creates a **resource group**, whose name uses the randomly-generated integer by a reference to the above resource

```
21  # Create the resource group
22
23      "ContactBookRG${random_integer.ri.result}"
24      "West Europe"
25
26
```

- Creates an **App Service Plan** with name, location (reference the location from the resource group), resource group name (reference the name of the resource group), operating system (set to "Linux") and type of SKU (set to "F1")

```
27  # Create the Linux App Service Plan
28
29      "contact-book-${random_integer.ri.result}"
30      azurerm_resource_group.rg.location
31      azurerm_resource_group.rg.name
32      "Linux"
33      "F1"
34
35
```

- Creates an **Azure Linux Web app** with **name**, **location**, **resource group name** and the **id of the service plan** (use **references** to the above resources)

```
36 # Create the web app, pass in the App Service Plan ID
37
38
39
```

```
40
41
42
```

- In addition, you should **add site configurations** including the **app's Node.js version** and a restriction for the **app to not always be on** (as we use the **free pricing plan**)

```
43 site_config {
44   application_stack {
45     node_version = "16-lts"
46   }
47   always_on = false
48 }
49
50
```

- **Deploys code** from the <https://github.com/nakov/ContactBook> repo, providing the **Web app id**, the **URL of the repo** and the **main branch name**

```
51 # Deploy code from a public GitHub repo
52
53
54
55
56 use_manual_integration = true
57
```

- Moreover, we should set the **use_manual_integration** argument to **true**, so that we **agree to deploy the app and its updates manually** when we use an **external Git** (a public GitHub repo, which is not our own and we cannot run CI/CD in GitHub Actions)

When **ready with the configuration file**, **initialize Terraform**, **format and validate the configuration** and **provision the resources** from the file. Know that this may **take a while**. It should be **successful** at the end:

```
Apply complete! Resources: 5 added, 0 changed, 0 destroyed.
```

When **done**, make sure that you **have a resource group**, an **app service plan** and a **Web app** in **Azure**:

Microsoft Azure

Search resources, services, and docs (G+)

Home > Resource groups > ContactBookRG55781

Resource group

Search

Create Manage view Delete resource group Refresh

Overview

Activity log

Access control (IAM)

Tags

Resource visualizer

Events

Essentials

JSON View

Resources Recommendations

| Name | Type | Location |
|-------------------------|------------------|-------------|
| contact-book-plan-55781 | App Service plan | West Europe |
| contact-book-55781 | App Service | West Europe |

Also, make sure that the "Contact Book" app is up and working on the provided domain URL in Azure. First, however, you should wait a bit and make sure that the deployment is successful:

Microsoft Azure

Search resources, services, and docs (G+)

Home > Resource groups > ContactBookRG55781 > contact-book-55781

Web App

Search

Browse Stop Swap Restart Delete Refresh

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Deployment Center

Deployment logs

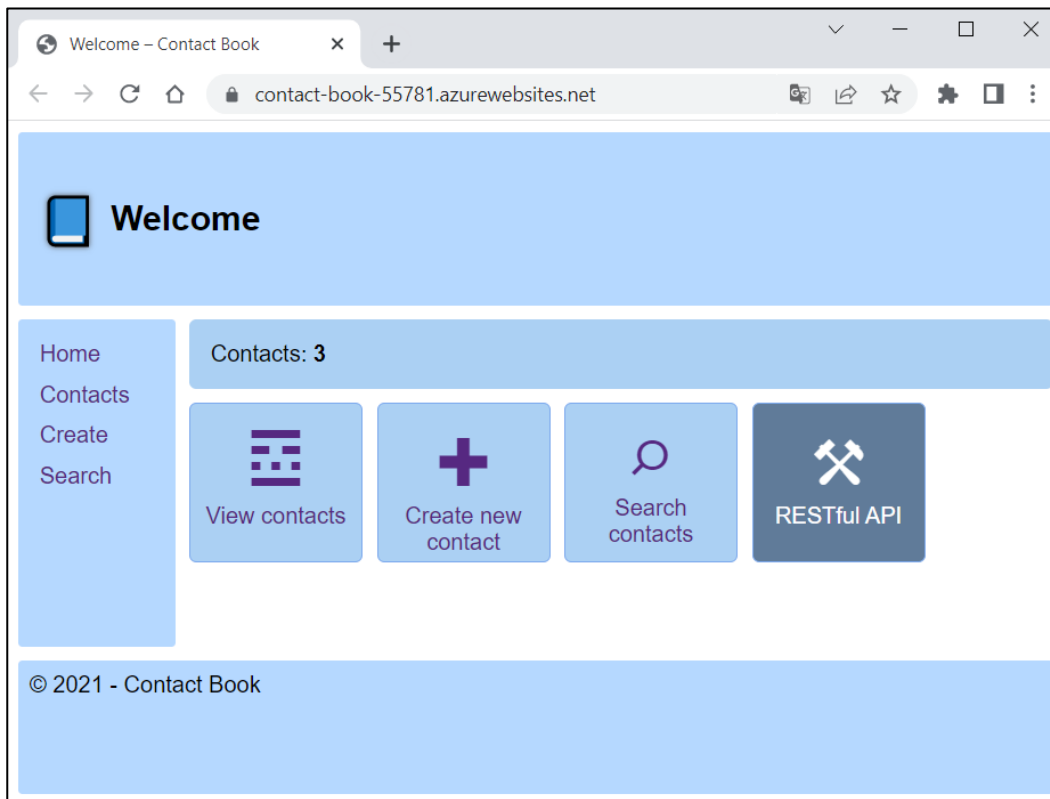
Last deployment

Deployment provider

View logs

Successful on Thursday, April 20, 01:58:12 PM Refresh

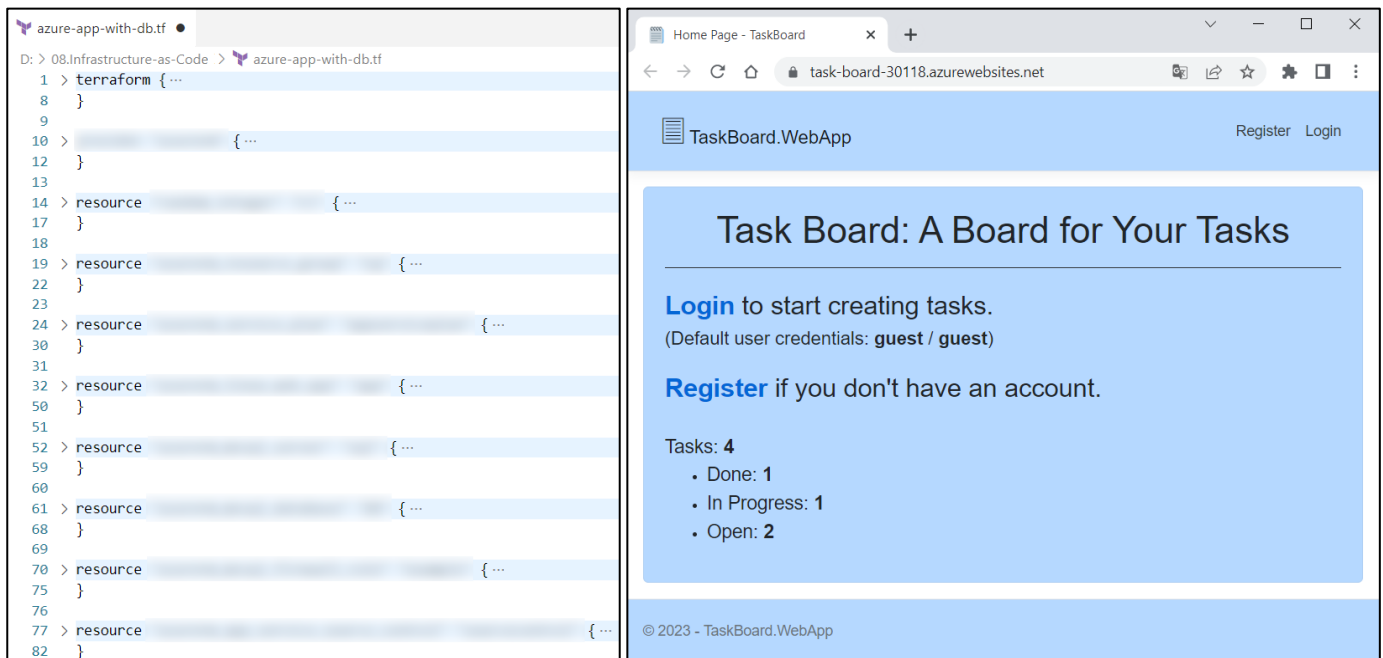
ExternalGit



Finally, you can **destroy the created Azure resources** using the **well-known Terraform command**.
And this is how you can **deploy an app to Azure** with some easy steps, using **Terraform**.

3. Azure Web App with Database

Create a **Terraform** configuration to create and deploy the "TaskBoard" Web app from the resources to **Azure Web Apps**. It is an **ASP.NET Core Web app** with a **SQL Server database**, which you should **upload to a GitHub repo** before you start.



Write and Apply a Terraform Configuration

In this task, you can use the **Terraform** configuration from the previous task but you should make the following **modifications and additions**:

- Create a **server resource in Azure** with **name**, **resource group name**, **location**, **version**, **administrator username** and **administrator password** arguments
- Create a **database resource in Azure** with **name**, **server ID**, **collation**, **license type**, **SKU name** and **zone redundancy** arguments
- Create a **firewall rule for the Azure server**, which has a **name** and **server ID** and sets **"0.0.0.0"** as **start and end IP addresses** (this means that it allows other **Azure resources** to access the server)
- **Application stack** should be set to **dotnet_version = "6.0"**
- The **Linux Web app** should contain a **connection_string** block with:
 - Name: **"DefaultConnection"**
 - Type: **"SQLAzure"**
 - Value: **"Data Source=tcp:\${fully qualified domain name of the MSSQL server},1433;Initial Catalog=\${name of the SQL database};User ID=\${username of the MSSQL server administrator};Password=\${password of the MSSQL server administrator};Trusted_Connection=False;MultipleActiveResultSets=True;"**
- The **GitHub repo URL** should be changed to point out a **repo with the source code** of the "TaskBoard" app

Find the **Terraform** resources you need and **how to configure them** by yourself. Also, use the **random integer** you have created as a resource to **generate unique names**, as well as **resource references** where possible.

When your **configuration is written**, use the well-known **Terraform** commands to **apply** it. After a while, your **declared resources should be provisioned in Azure**:

Microsoft Azure

Search resources, services, and docs (G+/)

Home >

All resources

Software University (SoftUni) (softwareuniversity.onmicrosoft.com)

+ Create Manage view Refresh Export to CSV Open query Assign tags Delete

Filter for any field... Subscription equals all Resource group equals all Add filter More (2)

0 Unsecure resources 0 Recommendations No grouping

List view

| Name | Type | Resource group | Location | Subscription |
|------------------------------------|------------------|------------------|--------------|--------------------|
| task-board-30118 | App Service | TaskBoardRG30118 | North Europe | Azure for Students |
| task-board-plan-30118 | App Service plan | taskboardrg30118 | North Europe | Azure for Students |
| task-board-sql-30118 | SQL server | TaskBoardRG30118 | North Europe | Azure for Students |
| TaskBoardDB30118 (task-board-sq... | SQL database | TaskBoardRG30118 | North Europe | Azure for Students |

And then, when the **app is deployed from the GitHub repo**, your **app should be up and working**.

Separate Configuration to Multiple Files

What we should do now is **separate our Terraform configuration to multiple files**, as it is **good practice** that allows **configuration modularity, reusability**, etc.

When done, we will have the **following files** (not necessary with the same file names):

- **main.tf** – the main Terraform configuration file
- **variables.tf** – contains variable declarations
- **values.tfvars** – contains values for the variables
- **outputs.tf** – contains outputs declarations

Let's see how to **separate our configuration**.

Step 1: Define Input Variables

You have the **configuration for provisioning and deploying a Web app with database** but it is all in one **.tf file** – including resource names, administrator credentials, etc. There are quite a **few hard coded values** that would make sense to have as **input parameters instead**, as this would allow us to **re-use the same template** to create multiple environments with a slightly different configuration.

In our **configuration**, we have the following **values** that can be turned into **input parameters**:

- Resource group name
- Resource group location
- App service plan name
- App service name
- SQL server name
- SQL database name
- SQL administrator login username
- SQL administrator password
- Firewall rule name

- GitHub repo URL

Create a **new .tf file** in the **Terraform configuration directory** and let's **define the input variables**. Each **variable** will have a **name**, **type** and **description**. In addition, it can have a **default value** that you can add if you want.

Define each variable from the above list in this way:

```
variables.tf X
D: > SoftUni > azure-app-deploy-asp-sql > variables.tf
1 variable "resource_group_name" {
2     type      = string
3     description = "Resource group name in Azure"
4 }
```

You can go on with the **rest of the variables' definition by yourself**, following the **syntax** shown. At the end, you should have **10 variables**:

```
variables.tf ●
D: > SoftUni > azure-app-deploy-asp-sql > variables.tf
1 > variable "resource_group_name" { ...
4 }
5
6 > variable "resource_group_location" { ...
9 }
10
11 > variable "app_service_plan_name" { ...
14 }
15
16 > variable "app_service_name" { ...
19 }
20
21 > variable "sql_server_name" { ...
24 }
25
26 > variable "sql_database_name" { ...
29 }
30
31 > variable "sql_admin_login" { ...
34 }
35
36 > variable "sql_admin_password" { ...
39 }
40
41 > variable "firewall_rule_name" { ...
44 }
45
46 > variable "repo_URL" { ...
49 }
```

Now let's **use these variables** in the **main Terraform configuration file** we have. To do this, use the following **syntax**: **var.{variable name}**. Do it like this for **all input variables** you defined:

```
main.tf x
D: > SoftUni > azure-app-deploy-asp-sql > main.tf
1 > terraform { ...
8 }
9
10 > provider "azurerm" { ...
12 }
13
14 resource "azurerm_resource_group" "rg" {
15     name      = var.resource_group_name
16     location = var.resource_group_location
17 }
```

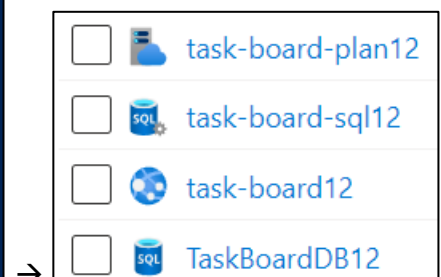
In addition, you can still use the **randomly generated integer value** as **part of the resource names** or you can **remove this resource** if you don't need it. However, make sure that your **resource names are unique enough** or **errors may appear**.

Now let's try to **apply the Terraform configuration** we have and see what will happen:

```
PS D:\SoftUni\azure-app-deploy-asp-sql> terraform apply
var.app_service_name
App Service name in Azure
Enter a value: _
```

As you can see, you are **prompted to enter an app service name** for the **app_service_name** input variable. You should **add values for all variables** and then they will be **used in your configuration**. All of them are **required** as we didn't put default values.

```
PS D:\SoftUni\azure-app-deploy-asp-sql> terraform apply
var.app_service_name
App Service name in Azure
Enter a value: task-board12
var.app_service_plan_name
App Service Plan name in Azure
Enter a value: task-board-plan12
...
var.sql_database_name
SQL Database name in Azure
Enter a value: TaskBoardDB12
var.sql_server_name
SQL Server instance name in Azure
Enter a value: task-board-sql12
```



Now we have **input variables for our configuration**, which is nice. However, if we run **terraform destroy**, we should **enter the same values again**, which is not pleasant.

Step 2: Create File with Variable Values

If we **don't want to enter values for the input variables**, we can **create a file** for them. Create a **file** with the **.tfvars** extension and **add value for each variable** using this syntax: **{variable name} = "{variable value}"**.

```
values.tfvars
D: > SoftUni > azure-app-deploy-asp-sql > values.tfvars
1  resource_group_name      = "TaskBoardRG12"
2  resource_group_location  = "North Europe"
3  app_service_plan_name    = "task-board-plan12"
4  app_service_name         = "task-board12"
5  sql_server_name          = "task-board-sql12"
6  sql_database_name        = "TaskBoardDB12"
7  sql_admin_login          = "user01"
8  sql_admin_password       = "@Aa123456789!"
9  firewall_rule_name       = "TaskBoardFirewallRule12"
10 repo_URL                 = "https://github.com/SoftUni-Projects/TaskBoard12"
```

Now we can **apply our configuration** again, using the **.tfvars** file we created:

```
PS D:\SoftUni\azure-app-deploy-asp-sql> terraform apply -var-file="values.tfvars"
```

The **file should be found** and **values used** – you should **not be prompted** to add any value manually.

Step 3: Define Outputs

At the end, we can **add outputs** that will **print us the URL of the Azure Web app** that will be created and its **outbound IP addresses**. **Outputs** are basically just pieces **state information** that you want to have available for different purposes.

You should create a **new .tf** file and **define the outputs** with **name** and **value** using the following syntax:

```
outputs.tf
D: > SoftUni > azure-app-deploy-asp-sql > outputs.tf
1  output "webapp_url" {
2    | value = azurerm_linux_web_app.app.default_hostname
3  }
4
5  output "webapp_ips" {
6    | value = azurerm_linux_web_app.app.outbound_ip_addresses
7  }
```

When you **apply the configuration**, the **values of the outputs** should be **printed in the terminal**:

```
PS D:\SoftUni\azure-app-deploy-asp-sql> terraform apply -var-file="values.tfvars"
...
Plan: 7 to add, 0 to change, 0 to destroy.

Changes to outputs:
+ webapp_ips = (known after apply)
+ webapp_url = (known after apply)

Do you want to perform these actions?
...
Apply complete! Resources: 7 added, 0 changed, 0 destroyed.

Outputs:
webapp_ips = "4.231.131.239,4.231.131.181,4.231.132.10,4.231.132.14,4.231.132.30,4.231.132.34,20.107.224.7"
webapp_url = "task-board12.azurewebsites.net"
```

After all this separation of the **Terraform configuration to files**, it should still be **working** and **provision the resources in Azure** successfully.

Now your **configuration follows good practices**. However, in the **next task** we will see how to **improve it** even more.





4. Terraform with CI/CD

Now we will **upload the Terraform configuration from the previous task** (for provisioning **Azure resources** and **deploying the "TaskBoard" Web app** to Azure Web Apps) to **GitHub** and will use **GitHub Actions workflows** to **test and run the configuration**.

By combining **Terraform with GitHub Actions**, we can **automate the infrastructure provisioning process**, ensure **consistency**, and **integrate it into your CI/CD workflows**, promoting **efficient software delivery** and **reducing manual tasks**. It provides a streamlined and efficient **workflow for managing infrastructure as code**, making it easier to **maintain, test, and deploy your infrastructure resources**.

We will have **GitHub Actions workflows** that will provision the **Azure resources** we want:

| Terraform Test | Terraform Plan | Terraform Apply |
|---------------------------|--------------------------------|-----------------------------------|
| succeeded now in 7s | succeeded 2 minutes ago in 40s | succeeded 2 minutes ago in 5m 12s |
| > ✓ Set up job 1s | > ✓ Set up job 5s | > ✓ Set up job 2s |
| > ✓ Checkout 0s | > ✓ Checkout 1s | > ✓ Checkout 0s |
| > ✓ Setup Terraform 1s | > ✓ Login via Azure CLI 18s | > ✓ Setup Terraform 1s |
| > ✓ Terraform Init 1s | > ✓ Setup Terraform 1s | > ✓ Terraform Init 1s |
| > ✓ Terraform Format 0s | > ✓ Terraform Init 2s | > ✓ Download Terraform Plan 0s |
| > ✓ Terraform Validate 0s | > ✓ Terraform Plan 11s | > ✓ Terraform Apply 5m 6s |
| > ✓ Post Checkout 0s | > ✓ Publish Terraform Plan 0s | > ✓ Post Checkout 0s |
| > ✓ Complete job 0s | > ✓ Post Checkout 0s | > ✓ Complete job 0s |
| > ✓ Complete job 0s | > ✓ Complete job 0s | > ✓ Complete job 0s |

| |
|--|
|  task-board-plan992244 |
|  task-board-sql992244 |
|  task-board992244 |
|  TaskBoardDB992244 (task-board-sql992244/TaskBoardDB992244) |

Start by creating a **GitHub repository**, which should contain your **main.tf** Terraform configuration file and your additional **Terraform files** – **terraform.tfvars** and **variables.tf**:

The screenshot shows the GitHub interface for a repository named 'Terraform-Actions-Azure'. The repository is public and has 1 branch (main) and 0 tags. It was created 7 minutes ago and has 2 commits. The file list shows: README.md, main.tf, terraform.tfvars, and variables.tf, all from the initial commit. The 'About' section indicates no description, website, or topics are provided. It also shows 0 stars, 1 watching, and 0 forks. The 'Releases' section shows no releases published.

Note: when the `.tfvars` file with variable values is named "terraform", Terraform finds it on its own and you should not point to it specifically in the Terraform commands you run.

Also, you **don't need the outputs .tf file**, as you can use GitHub Actions to show you what you need when a workflow is run.

Now let's see how to write the **GitHub Actions workflows** we need.

Test Workflow

We will first write a **test workflow in GitHub Actions** that will try to **initialize the working directory**, **check if the configuration files are correctly formatted** and **validate the configuration**.

Create a **YAML file in GitHub Actions**. The **workflow** should look like this:

The screenshot shows the GitHub interface for a workflow file named 'terraform-test.yml' in the repository 'Terraform-Actions-Azure'. The file is 32 lines long (26 loc) and 866 Bytes. The workflow is defined as follows:

```

1
2
3
4
5 jobs:
6   terraform-test:
7     # Checkout the repository to the GitHub Actions runner
8     # Run Terraform
9
10    steps:
11      # Checkout the repository to the GitHub Actions runner
12      # Run Terraform
13      # Run Terraform
14

```

```

15      # Install the latest version of the Terraform CLI
16
17
18
19
20
21      # Initialize a new or existing Terraform working directory
22      # Creating initial files, loading any remote state, downloading modules, etc.
23
24
25
26      # Checks that all Terraform configuration files adhere to a canonical format
27      - name: Terraform Format
28        run: terraform fmt -check -recursive
29
30      # Validate Terraform files
31
32

```

Look at the **comments in the above workflow** – they **describe the steps for testing the Terraform configuration**.

Write the workflow and **run it**. It should be **successful**:

The screenshot shows the GitHub Actions interface for a repository named 'Terraform-Actions-Azure'. The 'Actions' tab is selected, displaying a workflow named 'Create terraform-test.yml #1'. The workflow status is 'succeeded now in 7s'. The left sidebar shows the 'Jobs' section with 'Terraform Test' selected. The main panel displays the 'Terraform Test' job details, including a search bar for logs and a list of steps that all completed successfully:

| Step | Duration |
|--------------------|----------|
| Set up job | 1s |
| Checkout | 0s |
| Setup Terraform | 1s |
| Terraform Init | 1s |
| Terraform Format | 0s |
| Terraform Validate | 0s |
| Post Checkout | 0s |
| Complete job | 0s |

If you **receive any error**, **fix it** – you may have problems with your **Terraform configuration files** or the **workflow file** you have just created.

Apply Configuration Workflow

When we have a **valid configuration** with **working tests** in **GitHub Actions**, let's use a **workflow** to **provision resources** and **deploy the "TaskBoard" Web app** to **Azure**. You should **authenticate in Azure** using a **service principal** and then **write the workflow**.

Step 1: Create Service Principal
















We should **create a service principal** with a **"Contributor"** role in **Azure** that we will use to **authenticate GitHub Actions**. Do it with the **following command locally** or **manually through Azure Portal**:

```
PS C:\Users\PC> az ad sp create-for-rbac --name "Azure-Terraform-GitHub-Actions"
--role contributor --scopes /subscriptions/[redacted] --sdk-auth
Option '--sdk-auth' has been deprecated and will be removed in a future release.
Creating 'contributor' role assignment under scope '/subscriptions/8238f760-177c-42c1-ba0e-749ec718e461'
The output includes credentials that you must protect. Be sure that you do not include
these credentials in your code or check the credentials into your source control. For
more information, see https://aka.ms/azadsp-cli
{
  "clientId": "[redacted]",
  "clientSecret": "[redacted]",
  "subscriptionId": "[redacted]",
  "tenantId": "[redacted]",
  "activeDirectoryEndpointUrl": "https://login.microsoftonline.com",
  "resourceManagerEndpointUrl": "https://management.azure.com/",
  "activeDirectoryGraphResourceId": "https://graph.windows.net/",
  "sqlManagementEndpointUrl": "https://management.core.windows.net:8443/",
  "galleryEndpointUrl": "https://gallery.azure.com/",
  "managementEndpointUrl": "https://management.core.windows.net/"
}
```

Copy the **credentials JSON** as you will need it for the next step.

Step 2: Create GitHub Secrets

As you know, it is **good practice** to **store your credentials** as **secrets in GitHub**. You need the following secrets:

| Repository secrets | | |
|---|----------------------|---|
|  AZURE_CLIENT_ID | Updated 1 minute ago |   |
|  AZURE_CLIENT_SECRET | Updated 1 minute ago |   |
|  AZURE_CREDENTIALS | Updated 1 minute ago |   |
|  AZURE_SUBSCRIPTION_ID | Updated 1 minute ago |   |
|  AZURE_TENANT_ID | Updated 1 minute ago |   |

"AZURE_CREDENTIALS" should **contain the whole JSON** that we copied earlier and the **rest of the variables** should contain **only the corresponding parts** of it (only the **value**, without quotes "").

Now we are ready to write the **GitHub workflow** that uses these secrets.

Step 3: Write the Workflow

Finally, let's **write the workflow** that will consist of **2 jobs** – the first one will **create the Terraform plan** and the **second one will apply it**.

Write the workflow in this way:

```
Terraform-Actions-Azure / .github / workflows / terraform-plan-apply.yml
76 lines (62 loc) · 2.02 KB

Code Blame Raw Download Edit View Source

1
2
3
4
5 env:
6   ARM_CLIENT_ID: ${ secrets.AZURE_CLIENT_ID }
7   ARM_CLIENT_SECRET: ${ secrets.AZURE_CLIENT_SECRET }
8   ARM_SUBSCRIPTION_ID: ${ secrets.AZURE_SUBSCRIPTION_ID }
9   ARM_TENANT_ID: ${ secrets.AZURE_TENANT_ID }
10
11 jobs:
12   terraform-plan:
13     runs-on: ubuntu-latest
14     steps:
15       - name: Checkout the repository to the GitHub Actions runner
16         uses: actions/checkout@v2
17
18       - name: Login to Azure via Azure CLI
19         uses: azure/login@v1
20         with:
21           creds: ${ secrets.AZURE_CREDENTIALS }
22
23       - name: Install the latest version of the Terraform CLI
24         uses: hashicorp/setup-terraform@v1
25
26       - name: Initialize a new or existing Terraform working directory
27         # Creates initial files, loading any remote state, downloading modules, etc.
28         uses: terraform-docs/terraform-docs@v1
29
30       - name: Generates an execution plan for Terraform
31         uses: terraform-docs/terraform-docs@v1
32
33   terraform-apply:
34     runs-on: ubuntu-latest
35     needs: [terraform-plan]
36     steps:
37       - name: Checkout the repository to the GitHub Actions runner
38         uses: actions/checkout@v2
39
40       - name: Install the latest version of Terraform CLI
41         uses: hashicorp/setup-terraform@v1
42
43       - name: Initialize a new or existing Terraform working directory
44         # Creates initial files, loading any remote state, downloading modules, etc.
45         uses: terraform-docs/terraform-docs@v1
46
47       - name: Download saved plan from artifacts
48         uses: actions/download-artifact@v2
49
50       - name: Terraform Apply
51         uses: terraform-docs/terraform-docs@v1
52         with:
53           -auto-approve tfplan
```

You can use the **steps from the test workflow** we created earlier as part of **this YAML file**.

Note some **specific things** about this **workflow**:

- You need some **environment variables** so that **Terraform can authenticate in Azure**.
- You should use the **"AZURE_CREDENTIALS" GitHub secret** to **authenticate GitHub Actions in Azure**.
- The **second job** should **depend on the execution** of the **first one**.
- You should **add the "-auto-approve tfplan" flag** to **automatically approve the changes** in the **"tfplan"** without requiring manual confirmation during the workflow run.

The **workflow should run successfully**:

GitHub interface showing the workflow `terraform-plan-apply.yml` for the repository `Terraform-Actions-Azure`. The workflow is titled "Create terraform-plan-apply.yml".

The workflow summary shows two jobs: "Terraform Plan" and "Terraform Apply". The "Terraform Plan" job is highlighted.

The "Terraform Plan" job details show the following steps:

- Set up job (5s)
- Checkout (1s)
- Login via Azure CLI (18s)
- Setup Terraform (1s)
- Terraform Init (2s)
- Terraform Plan (11s)
- Publish Terraform Plan (0s)
- Post Checkout (0s)
- Complete job (0s)

GitHub interface showing the workflow `terraform-plan-apply.yml` for the repository `Terraform-Actions-Azure`. The workflow is titled "Create terraform-plan-apply.yml".

The workflow summary shows two jobs: "Terraform Plan" and "Terraform Apply". The "Terraform Apply" job is highlighted.

The "Terraform Apply" job details show the following steps:

- Set up job (2s)
- Checkout (0s)
- Setup Terraform (1s)
- Terraform Init (1s)
- Download Terraform Plan (0s)
- Terraform Apply (5m 6s)
- Post Checkout (0s)
- Complete job (0s)

Also, the **Azure resources** you defined in the **Terraform configuration** should be **provisioned** and the **"TaskBoard"** app deployed and working:

Microsoft Azure Search resources, services, and docs (G+/)

Home >

All resources

Software University (SoftUni) (softwareuniversity.onmicrosoft.com)

+ Create Manage view Refresh Export to CSV Open query Assign tags Delete

Filter for any field... Subscription equals all Add filter More (3)

0 Unsecure resources 0 Recommendations

No grouping List view

| Name | Type | Resource group | Location | Subscription |
|-------------------------|------------------|-------------------|--------------|--------------------|
| task-board-plan992244 | App Service plan | TaskBoardRG992244 | North Europe | Azure for Students |
| task-board-sql992244 | SQL server | TaskBoardRG992244 | North Europe | Azure for Students |
| task-board992244 | App Service | TaskBoardRG992244 | North Europe | Azure for Students |
| TaskBoardDB992244 (...) | SQL database | TaskBoardRG992244 | North Europe | Azure for Students |

Page 1 of 1 Showing 1 to 4 of 4 records. Give feedback

Home Page - TaskBoard

task-board992244.azurewebsites.net

TaskBoard.WebApp Register Login

Task Board: A Board for Your Tasks

Login to start creating tasks.
(Default user credentials: **guest** / **guest**)

Register if you don't have an account.

Tasks: 4

- Done: 1
- In Progress: 1
- Open: 2

© 2023 - TaskBoard.WebApp

We successfully used **GitHub Actions** to run a **Terraform** configuration that **provisions resources in Azure**. However, if we **change the configuration** and **run the workflow again**, an **error will occur**. This happens because we **don't save the Terraform configuration state file**.

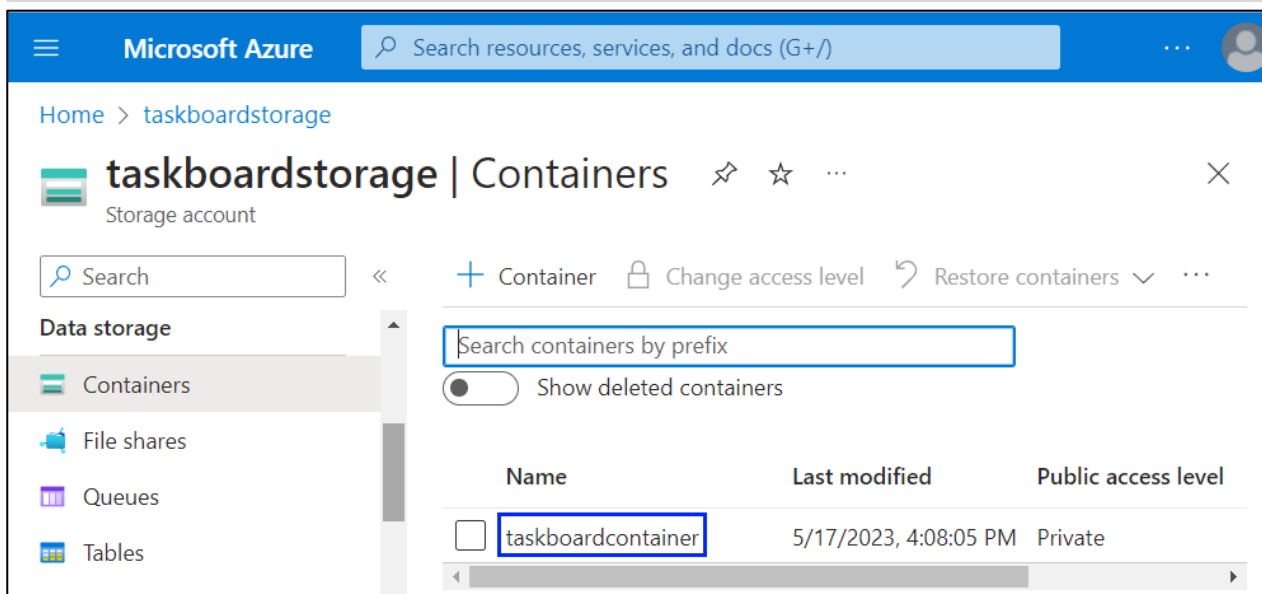
Store State File in Azure Storage Account

Terraform utilizes a **state file** to **store information** about the **current state of your managed infrastructure** and associated configuration. This file will need to be **persisted between different runs of the workflow**.

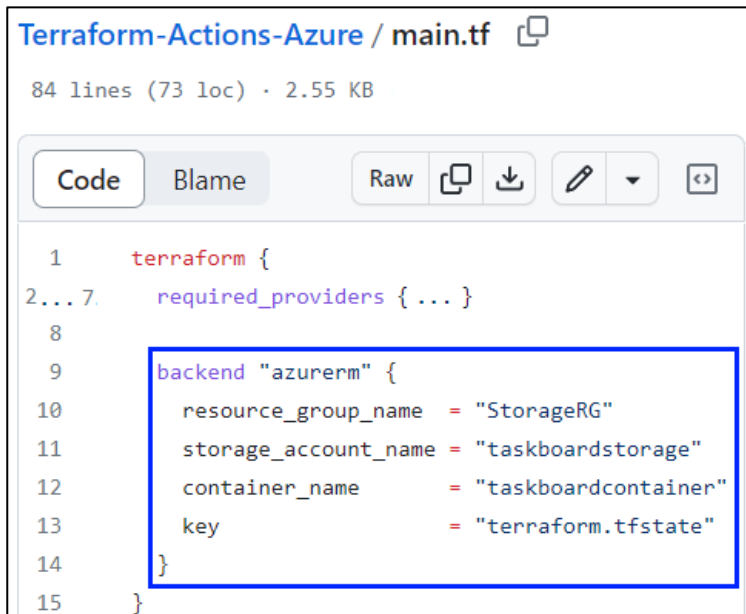
The recommended approach is to **store this file** within an **Azure Storage Account** and this is what we will do now. First, you should **create an Azure storage account** with a **container** to **store the state file**:

```
PS C:\Users\PC> az storage account create
>> --name taskboardstorage
>> --resource-group StorageRG
>> --location northeurope
>> --sku Standard_LRS
>> --kind StorageV2
```

```
PS C:\Users\PC> az storage container create -n taskboardcontainer --account-name taskboardstorage
```



Then, to **use this storage in Terraform**, you should **add a backend block** in the **main.tf** configuration file:

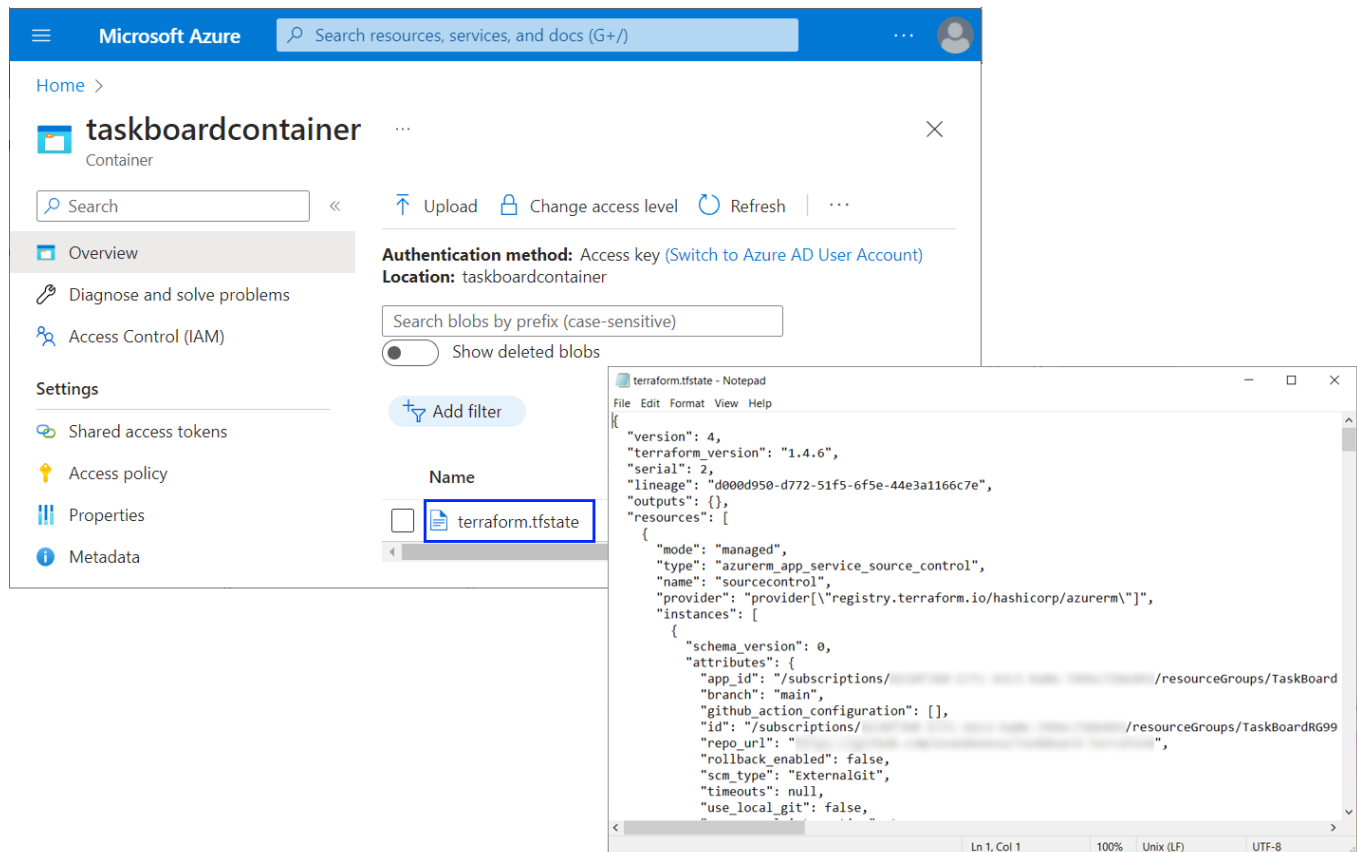


A **backend block** defines where **Terraform** stores its **state data files**. You should provide the **names of your resource group, storage account and container**, as well as to set a **name of the state file** that will be created.

Commit the changed file to GitHub and wait for **GitHub Actions** to run the workflow.

Note: you **should not have your resources in Azure now** or the **GitHub Actions workflow** will still give you an error when **trying to create them**, as they are **not defined in the state file**. **Delete the resources** you created previously with your Terraform configuration from **Azure**.

The **workflow should be successful** and you should see that a **terraform.tfstate** file was created in your **Azure storage container**:



Go and **make a change in your Terraform configuration in GitHub** and run the **workflow again** – the **modified resources** should be **updated successfully in Azure**.

Now you have a **fully working GitHub Actions + Terraform + Azure configuration** to create and manage resources.

* More Configuration Improvements

We have a **good Terraform configuration** and **GitHub Actions workflows** created during the previous tasks but here are some **additional challenges for you** to overcome to **improve your Terraform skills** even more:


- You can **create a Terraform configuration file to provision an Azure storage account and container** for the **Terraform backend**, instead of doing it with commands like we did previously. Then, you can use a **GitHub Actions workflow** to run that configuration and **provision the resources in Azure**.
- You can **create a Terraform configuration file to create the service principal and assign the "Contributor" role to it** instead of doing it manually or with commands through Azure CLI. You can again try to **run the configuration in GitHub Actions**, not only locally.

By **fulfilling these additional tasks**, you would have **fully explored** and used the **integration between Terraform, GitHub Actions and Azure**.

II. App Monitoring

5. Monitor the "Contact Book" Node.js App with Prometheus

We have the **Node.js "Contact Book" app** in the **resources**. We aim to **monitor it using Prometheus**, so we need its **metrics**. In this case, we will **instrument the app to expose the metrics** we want. And then we will **configure Prometheus to display these metrics**.




Step 1: Examine the App

We have the **"Contact Book" Node.js app**, which holds a **searchable list of contacts**. You have pages to **list all contacts** (`/contacts`), **view a single contact** (`/contacts/:id`), **search for a contact** (`/contacts/search/:keyword`) and **add a new contact** (`/contacts/create`).

Open the **project in Visual Studio Code** to **examine its files**:

Open a **terminal** and **execute the "npm install" and "npm start" commands to run the app**:



```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
node + - [ ] [ ] ... ^ X

PS D:\SoftUni\ContactBook> npm install
npm WARN deprecated core-js@2.6.12: core-js@<3.23.3 is no longer maintained and no
t recommended for usage due to the number of issues. Because of the V8 engine whim
s, feature detection in old core-js versions could cause a slowdown up to 100x eve
n if nothing is polyfilled. Some versions have web compatibility issues. Please, u
pgrade your dependencies to the actual version of core-js.

added 123 packages, and audited 124 packages in 7s

13 packages are looking for funding
  run `npm fund` for details
1 high severity vulnerability

To address all issues (including breaking changes), run:
  npm audit fix --force
Run `npm audit` for details.
```

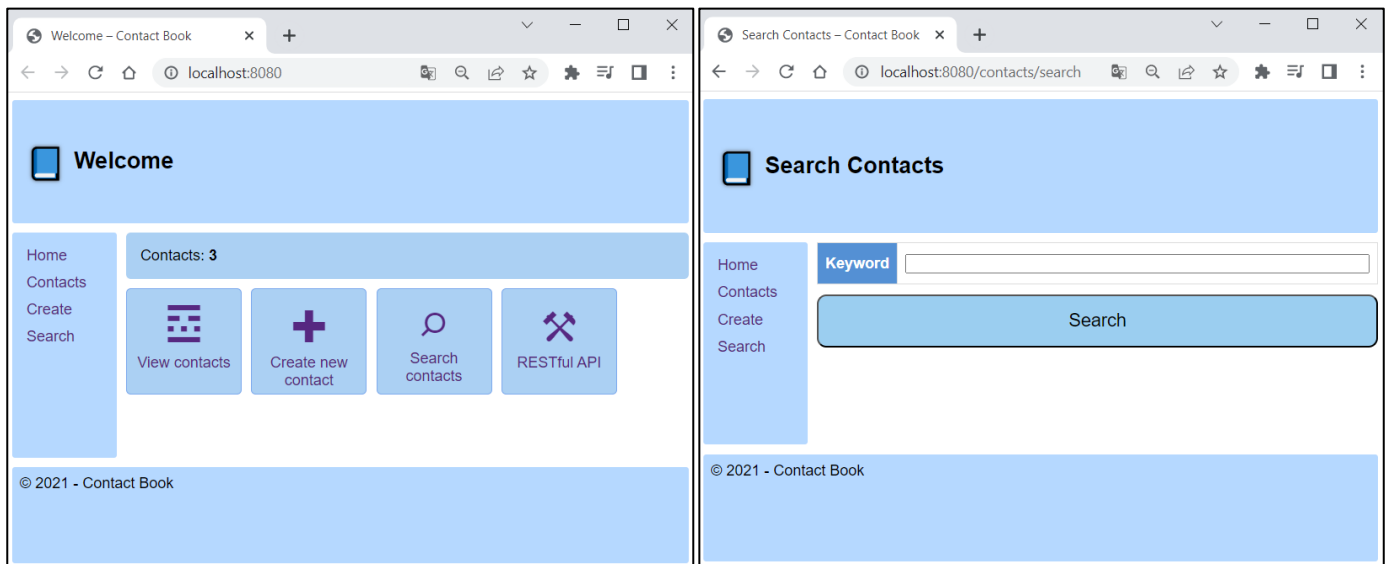
```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL
node + - [ ] [ ] ... ^ X

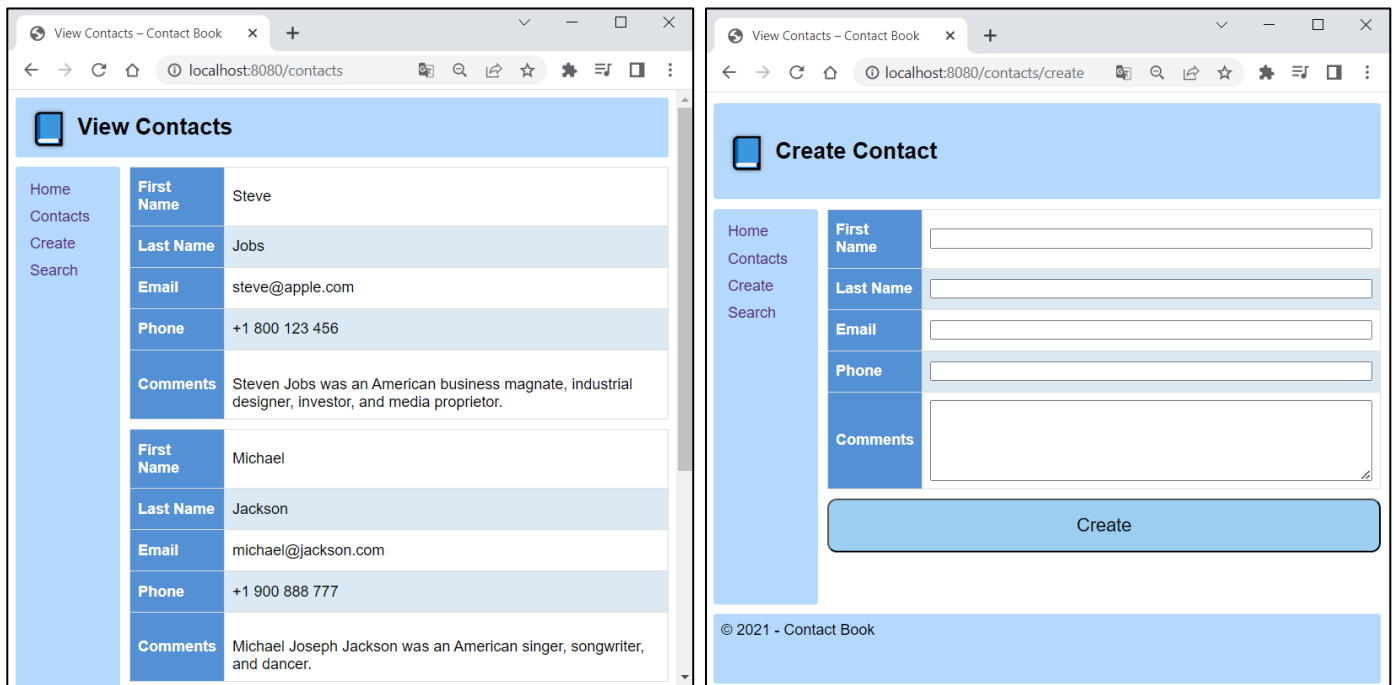
PS D:\SoftUni\ContactBook> npm start

> start
> node index.js

App started. Listening at http://localhost:8080
[ ]
```

Look at the **app pages** on <http://localhost:8080>:





Let's now see how to **modify the app code to export app metrics for Prometheus**.

Step 2: Export Node.js App Metrics

To make app metrics readable for Prometheus, we should install an additional client library for Node.js and then modify the code to define and export the metrics we want.

Install Prom-Client

Stop the app with [Ctrl] + [C] in the terminal. Then, we should install the **prom-client** package, which is the Prometheus client for Node.js that supports histogram, summaries, gauges and counters. Do it with the following command:

```

PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL
PS D:\SoftUni\ContactBook> npm install prom-client
added 3 packages, and audited 127 packages in 767ms

13 packages are looking for funding
  run `npm fund` for details

1 high severity vulnerability

To address all issues (including breaking changes), run:
  npm audit fix --force

Run `npm audit` for details.
  
```

You can see that the **prom-client** package is added to the **package.json** file with **project dependencies**:



```
{} package.json 1, M
{} package.json > ...
1 {
2   "name": "ContactBook",
3   "main": "index.js",
4   "dependencies": {
5     "body-parser": "^1.19.0",
6     "express": "^4.17.1",
7     "prom-client": "^14.2.0",
8     "pug": "^2.0.4"
9   },
10  "scripts": {
11    "start": "node index.js"
12  }
13 }
```

Because of this, you **won't need to install the package** separately from the others next time.

Export Default Metrics

Now we will **modify our code** to **collect the default app metrics** together with **some custom ones** and **expose them** on the **/metrics** endpoint.

To do this, navigate to **mvc-controller.js** file where the main app routing is and **include the prom-client module**, as we will need it:

```
JS mvc-controller.js
controllers > JS mvc-controller.js > setup
1 const client = require('prom-client');
2
```

Then, **create a registry** to register the metrics:

```
3 const register = new client.Registry();
4
```

Use the **collectDefaultMetrics()** function from the **imported module** to **collect and register default metrics** for **monitoring the Node.js application**, for example CPU usage, memory usage, event loop latency, and garbage collection duration:

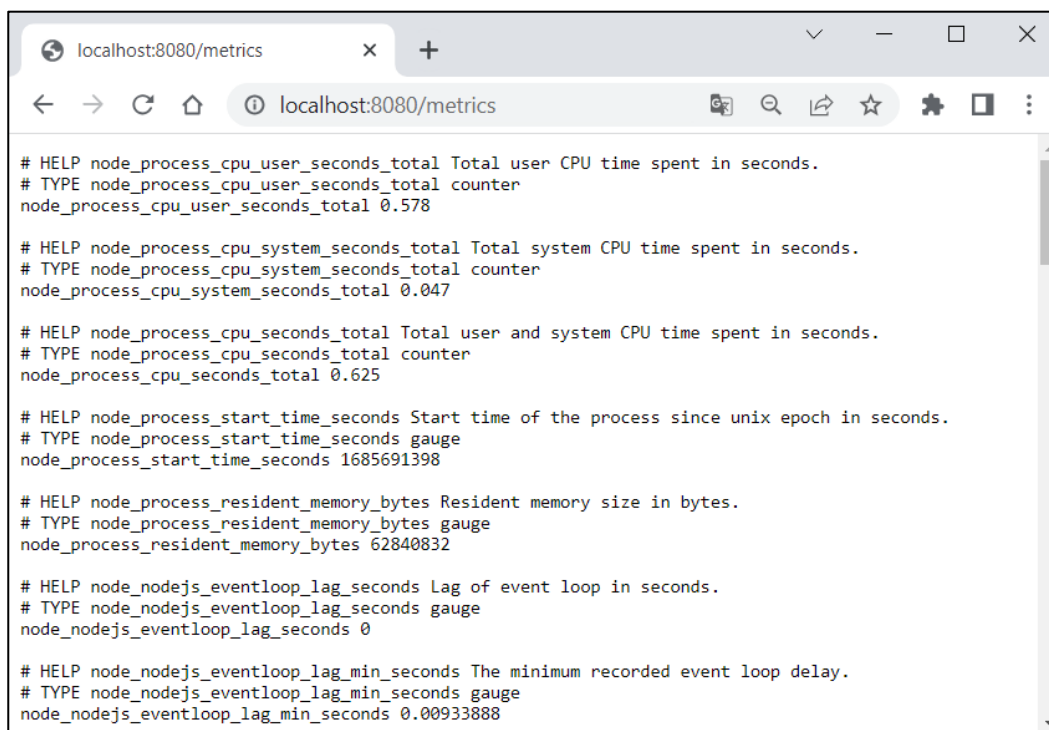
```
5 client.collectDefaultMetrics({
6   app: 'node-application-monitoring-app',
7   prefix: 'node_',
8   timeout: 10000,
9   gcDurationBuckets: [0.001, 0.01, 0.1, 1, 2, 5],
10  register
11 });
```

This configuration sets **default metric names** to start with the **"_node"** prefix, the **timeout** to 10000ms, the **buckets for the default metric that measures garbage collection (GC) durations** (values represent the upper bounds of each bucket) and the **registry** that we created to be used.

These are the **default metrics** we will export. Now, in the `setup()` function, set up an **HTTP GET route** for the **/metrics endpoint**, which should **return the collected app metrics as response**:

```
13 function setup(app, data) {
14   app.get('/metrics', async (req, res) => {
15     res.setHeader('Content-Type', register.contentType);
16     res.send(await register.metrics());
17   });
18 }
```

Before we add some **custom metrics**, let's see how **default metrics are showed**. Save the changes and **start the app** again. Then, **navigate to <http://localhost:8080/metrics>** in the browser:

A screenshot of a web browser window with the address bar showing 'localhost:8080/metrics'. The page content displays a list of Node.js default metrics in a text-based format. Each metric entry includes a help text, the metric type, and the current value. The metrics shown are: node_process_cpu_user_seconds_total (0.578), node_process_cpu_system_seconds_total (0.047), node_process_cpu_seconds_total (0.625), node_process_start_time_seconds (1685691398), node_process_resident_memory_bytes (62840832), node_nodejs_eventloop_lag_seconds (0), and node_nodejs_eventloop_lag_min_seconds (0.00933888).

```
# HELP node_process_cpu_user_seconds_total Total user CPU time spent in seconds.
# TYPE node_process_cpu_user_seconds_total counter
node_process_cpu_user_seconds_total 0.578

# HELP node_process_cpu_system_seconds_total Total system CPU time spent in seconds.
# TYPE node_process_cpu_system_seconds_total counter
node_process_cpu_system_seconds_total 0.047

# HELP node_process_cpu_seconds_total Total user and system CPU time spent in seconds.
# TYPE node_process_cpu_seconds_total counter
node_process_cpu_seconds_total 0.625

# HELP node_process_start_time_seconds Start time of the process since unix epoch in seconds.
# TYPE node_process_start_time_seconds gauge
node_process_start_time_seconds 1685691398

# HELP node_process_resident_memory_bytes Resident memory size in bytes.
# TYPE node_process_resident_memory_bytes gauge
node_process_resident_memory_bytes 62840832

# HELP node_nodejs_eventloop_lag_seconds Lag of event loop in seconds.
# TYPE node_nodejs_eventloop_lag_seconds gauge
node_nodejs_eventloop_lag_seconds 0

# HELP node_nodejs_eventloop_lag_min_seconds The minimum recorded event loop delay.
# TYPE node_nodejs_eventloop_lag_min_seconds gauge
node_nodejs_eventloop_lag_min_seconds 0.00933888
```

Now let's **add some more metrics**.

Export Custom Metrics

The **custom metrics** we shall export are about the **duration of HTTP requests to different endpoints in seconds**. They will be saved in a **histogram** with **buckets from 0.01 to 1 seconds** and will keep **request method, route and status code**.

Add the following code to create the histogram metric (before the `setup()` function):

```
13 const httpRequestTimer = new client.Histogram({
14   name: 'http_request_duration_seconds',
15   help: 'Duration of HTTP requests in seconds',
16   labelNames: ['method', 'route', 'code'],
17   buckets: [0.01, 0.03, 0.05, 0.07, 0.1, 0.3, 0.5, 0.7, 1]
18 });
19
```

Then we should **register the metric**:

```
20 register.registerMetric(httpRequestTimer);
21
```

Now, for **each of the routes**, we should:

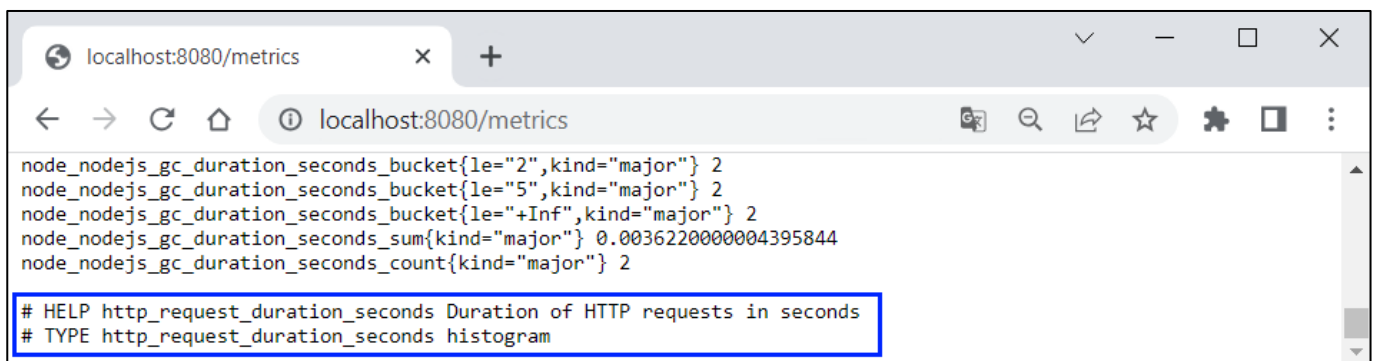
- Start an **HTTP request timer**, saving a **reference** to the returned method
- Save **reference to the path** so we can record it when ending the timer
- And finally **end the timer** and **add labels**

In this way, the **HTTP request data and duration** will be recorded. Do it for the **/metrics** endpoint like this:

```
22 function setup(app, data) {
23   app.get('/metrics', async (req, res) => {
24     const end = httpRequestTimer.startTimer();
25     const route = req.route.path;
26
27     res.setHeader('Content-Type', register.contentType);
28     res.send(await register.metrics());
29
30     end({ route, code: res.statusCode, method: req.method });
31   });
32 }
```

Do it for the **rest of the endpoint methods** in the same way by **adding the above three lines**. When ready, **run the app** again.

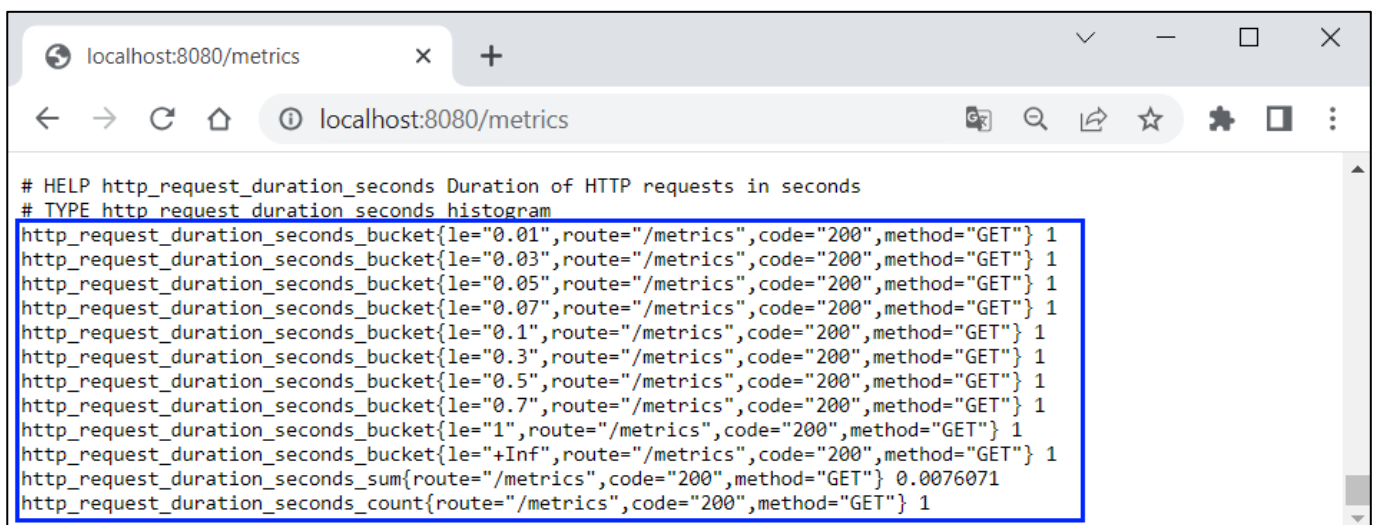
When you first **access /metrics**, you will see the **new metrics at the bottom**:



```
node_nodejs_gc_duration_seconds_bucket{le="2",kind="major"} 2
node_nodejs_gc_duration_seconds_bucket{le="5",kind="major"} 2
node_nodejs_gc_duration_seconds_bucket{le="+Inf",kind="major"} 2
node_nodejs_gc_duration_seconds_sum{kind="major"} 0.003622000004395844
node_nodejs_gc_duration_seconds_count{kind="major"} 2

# HELP http_request_duration_seconds Duration of HTTP requests in seconds
# TYPE http_request_duration_seconds histogram
```

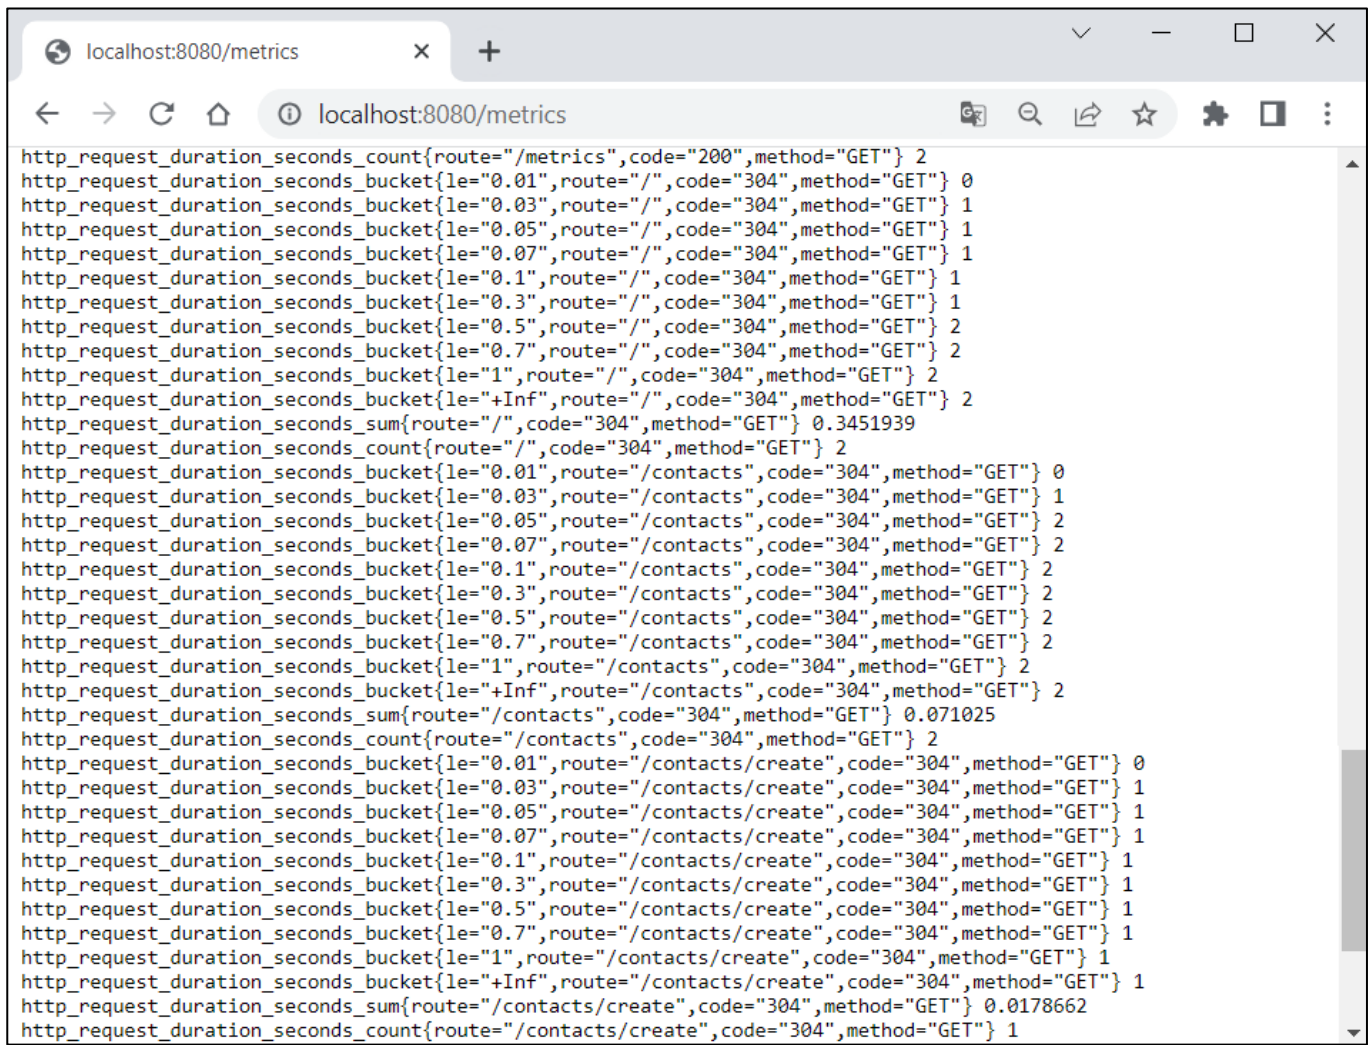
However, you still have **no metric values**. You should **refresh the page**, so that the metrics for the **previous HTTP request** to **/metrics** are displayed:



```
# HELP http_request_duration_seconds Duration of HTTP requests in seconds
# TYPE http_request_duration_seconds histogram
http_request_duration_seconds_bucket{le="0.01",route="/metrics",code="200",method="GET"} 1
http_request_duration_seconds_bucket{le="0.03",route="/metrics",code="200",method="GET"} 1
http_request_duration_seconds_bucket{le="0.05",route="/metrics",code="200",method="GET"} 1
http_request_duration_seconds_bucket{le="0.07",route="/metrics",code="200",method="GET"} 1
http_request_duration_seconds_bucket{le="0.1",route="/metrics",code="200",method="GET"} 1
http_request_duration_seconds_bucket{le="0.3",route="/metrics",code="200",method="GET"} 1
http_request_duration_seconds_bucket{le="0.5",route="/metrics",code="200",method="GET"} 1
http_request_duration_seconds_bucket{le="0.7",route="/metrics",code="200",method="GET"} 1
http_request_duration_seconds_bucket{le="1",route="/metrics",code="200",method="GET"} 1
http_request_duration_seconds_bucket{le="+Inf",route="/metrics",code="200",method="GET"} 1
http_request_duration_seconds_sum{route="/metrics",code="200",method="GET"} 0.0076071
http_request_duration_seconds_count{route="/metrics",code="200",method="GET"} 1
```

As you can see, the **first HTTP request** to **/metrics** took about **0.0076 seconds**, which is **less than 0.01** and that's why it falls into **each of the buckets**.

If you **access the other app endpoints**, you will get even **more metric data**:




```
http_request_duration_seconds_count{route="/metrics",code="200",method="GET"} 2
http_request_duration_seconds_bucket{le="0.01",route="/",code="304",method="GET"} 0
http_request_duration_seconds_bucket{le="0.03",route="/",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.05",route="/",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.07",route="/",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.1",route="/",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.3",route="/",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.5",route="/",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="0.7",route="/",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="1",route="/",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="+Inf",route="/",code="304",method="GET"} 2
http_request_duration_seconds_sum{route="/",code="304",method="GET"} 0.3451939
http_request_duration_seconds_count{route="/",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="0.01",route="/contacts",code="304",method="GET"} 0
http_request_duration_seconds_bucket{le="0.03",route="/contacts",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.05",route="/contacts",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="0.07",route="/contacts",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="0.1",route="/contacts",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="0.3",route="/contacts",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="0.5",route="/contacts",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="0.7",route="/contacts",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="1",route="/contacts",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="+Inf",route="/contacts",code="304",method="GET"} 2
http_request_duration_seconds_sum{route="/contacts",code="304",method="GET"} 0.071025
http_request_duration_seconds_count{route="/contacts",code="304",method="GET"} 2
http_request_duration_seconds_bucket{le="0.01",route="/contacts/create",code="304",method="GET"} 0
http_request_duration_seconds_bucket{le="0.03",route="/contacts/create",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.05",route="/contacts/create",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.07",route="/contacts/create",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.1",route="/contacts/create",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.3",route="/contacts/create",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.5",route="/contacts/create",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="0.7",route="/contacts/create",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="1",route="/contacts/create",code="304",method="GET"} 1
http_request_duration_seconds_bucket{le="+Inf",route="/contacts/create",code="304",method="GET"} 1
http_request_duration_seconds_sum{route="/contacts/create",code="304",method="GET"} 0.0178662
http_request_duration_seconds_count{route="/contacts/create",code="304",method="GET"} 1
```

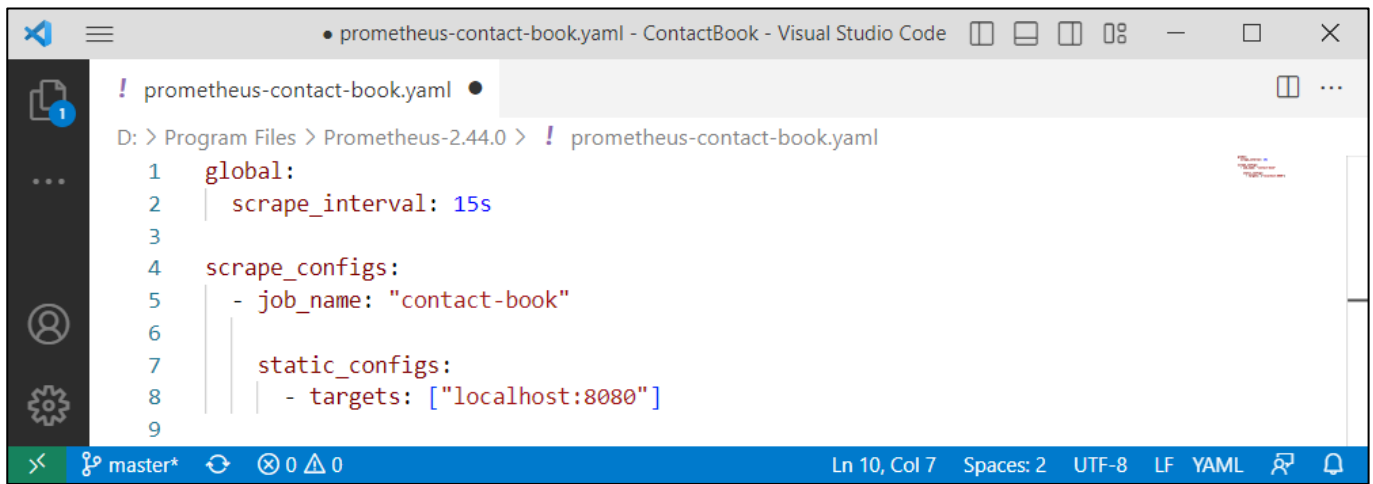
You can use the **browser inspector** in the browser to **compare the HTTP request times** shown here and there – they should be pretty close as values.

Step 3: Condifure and Run Prometheus

Go to the **Prometheus installation directory** where our binary files are and **create a YAML file** where we will **write the configuration** for monitoring the "Contact Book" app:

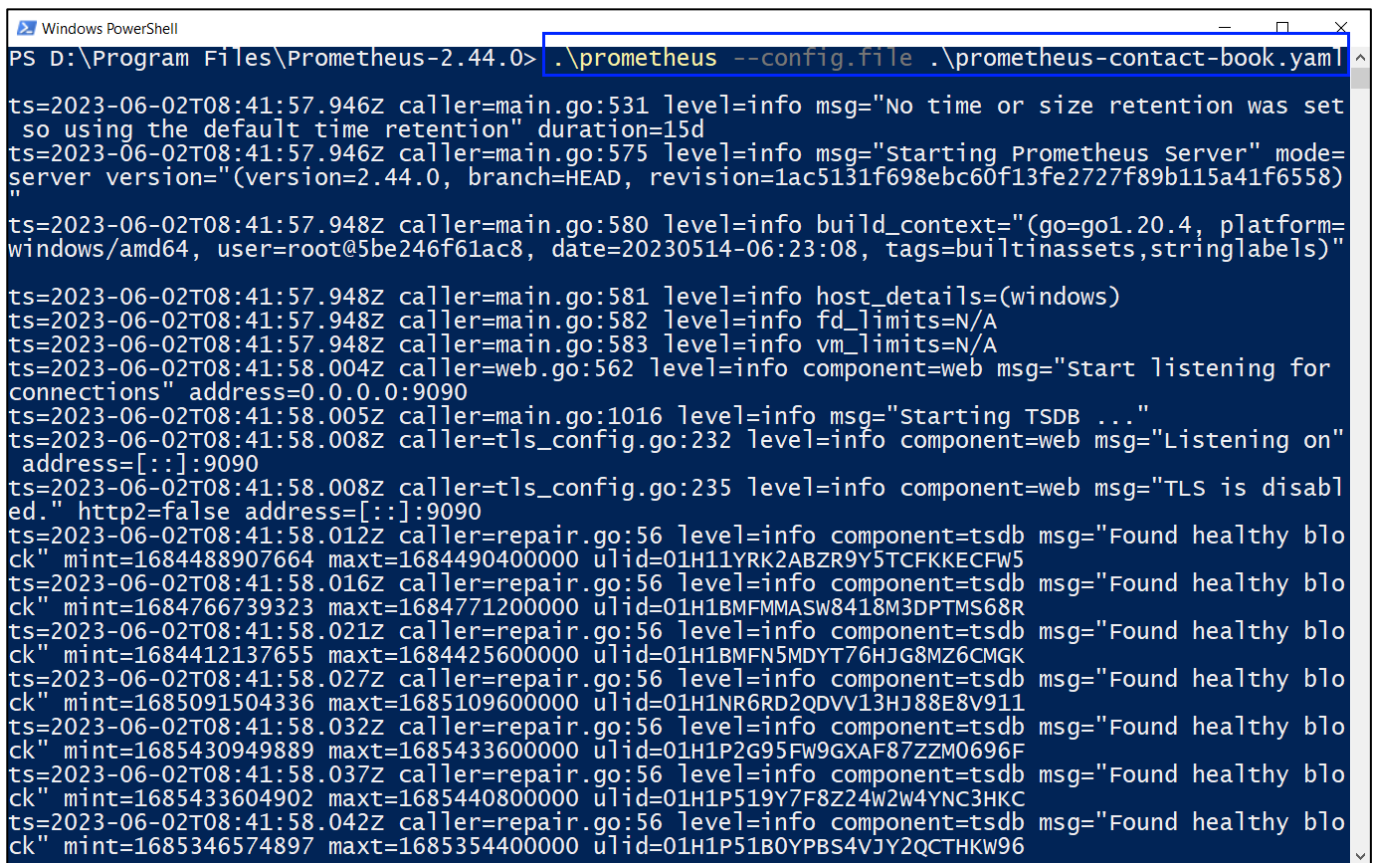


In the **Prometheus configuration file**, we should **define a single job** to monitor our app on <http://localhost:8080> and **scrape target metrics** on **every 15 seconds**:



```
1 global:
2   scrape_interval: 15s
3
4 scrape_configs:
5   - job_name: "contact-book"
6
7     static_configs:
8       - targets: ["localhost:8080"]
9
```

Save the file and open a terminal. Navigate to the **Prometheus installation directory** and run **Prometheus** with this configuration file:



```
PS D:\Program Files\Prometheus-2.44.0> .\prometheus --config.file .\prometheus-contact-book.yaml
ts=2023-06-02T08:41:57.946Z caller=main.go:531 level=info msg="No time or size retention was set so using the default time retention" duration=15d
ts=2023-06-02T08:41:57.948Z caller=main.go:575 level=info msg="Starting Prometheus server" mode=server version="(version=2.44.0, branch=HEAD, revision=1ac5131f698ebc60f13fe2727f89b115a41f6558)"
ts=2023-06-02T08:41:57.948Z caller=main.go:580 level=info build_context="(go=go1.20.4, platform=windows/amd64, user=root@5be246f61ac8, date=20230514-06:23:08, tags=builtinassets,stringlabels)"
ts=2023-06-02T08:41:57.948Z caller=main.go:581 level=info host_details=(windows)
ts=2023-06-02T08:41:57.948Z caller=main.go:582 level=info fd_limits=N/A
ts=2023-06-02T08:41:57.948Z caller=main.go:583 level=info vm_limits=N/A
ts=2023-06-02T08:41:58.004Z caller=web.go:562 level=info component=web msg="Start listening for connections" address=0.0.0.0:9090
ts=2023-06-02T08:41:58.005Z caller=main.go:1016 level=info msg="Starting TSDB ..."
ts=2023-06-02T08:41:58.008Z caller=tsdb.go:232 level=info component=web msg="Listening on" address=[::]:9090
ts=2023-06-02T08:41:58.008Z caller=tsdb.go:235 level=info component=web msg="TLS is disabled." http2=false address=[::]:9090
ts=2023-06-02T08:41:58.012Z caller=repair.go:56 level=info component=tsdb msg="Found healthy block" mint=1684488907664 maxt=1684490400000 ulid=01H11YRK2ABZR9Y5TCFKKECFW5
ts=2023-06-02T08:41:58.016Z caller=repair.go:56 level=info component=tsdb msg="Found healthy block" mint=1684766739323 maxt=1684771200000 ulid=01H1BMFMASW8418M3DPTMS68R
ts=2023-06-02T08:41:58.021Z caller=repair.go:56 level=info component=tsdb msg="Found healthy block" mint=1684412137655 maxt=1684425600000 ulid=01H1BMFN5MDYT76HJG8MZ6CMGK
ts=2023-06-02T08:41:58.027Z caller=repair.go:56 level=info component=tsdb msg="Found healthy block" mint=1685091504336 maxt=1685109600000 ulid=01H1NR6RD2QDVV13HJ88E8V911
ts=2023-06-02T08:41:58.032Z caller=repair.go:56 level=info component=tsdb msg="Found healthy block" mint=1685430949889 maxt=1685433600000 ulid=01H1P2G95FW9GXA87ZZM0696F
ts=2023-06-02T08:41:58.037Z caller=repair.go:56 level=info component=tsdb msg="Found healthy block" mint=1685433604902 maxt=1685440800000 ulid=01H1P519Y7F8Z24W2W4YNC3HKC
ts=2023-06-02T08:41:58.042Z caller=repair.go:56 level=info component=tsdb msg="Found healthy block" mint=1685346574897 maxt=1685354400000 ulid=01H1P51B0YPBS4VJY2QCTHKW96
```

Prometheus server should be available on <http://localhost:9090> by default:

Prometheus Time Series Collectio x +

localhost:9090/graph?g0.expr=&g0.tab=...

Prometheus Alerts Graph Status Help

☒ Use local time ☐ Enable query history ☒ Enable autocomplete ☒ Enable highlighting ☒ Enable linter

Search Expression (press Shift+Enter for newlines) Execute

Table Graph

Evaluation time

No data queried yet

Remove Panel

Add Panel

You can navigate to **[Status] → [Targets]** to see that **Prometheus connects to the configured target app successfully**:

Prometheus Alerts Graph Status Help

Targets

All scrape pools All Unhealthy Collapse All

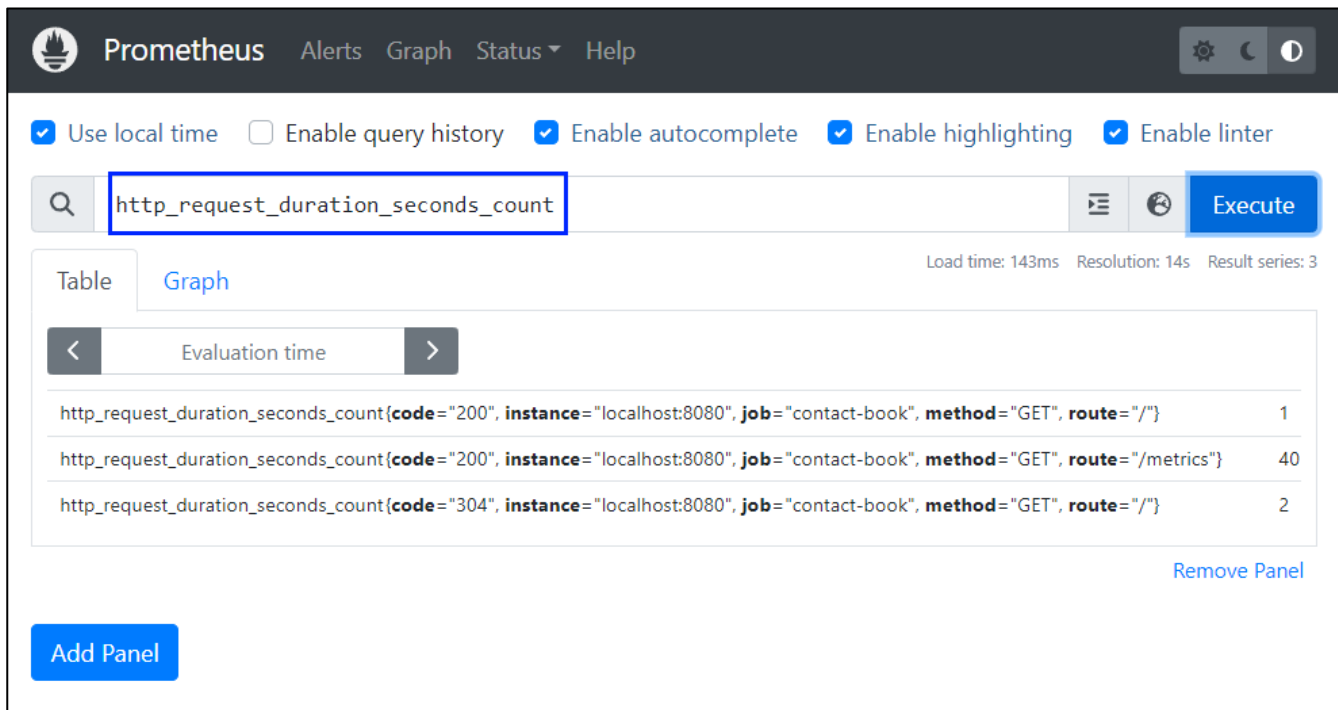
Unknown Unhealthy Healthy

contact-book (1/1 up) show less

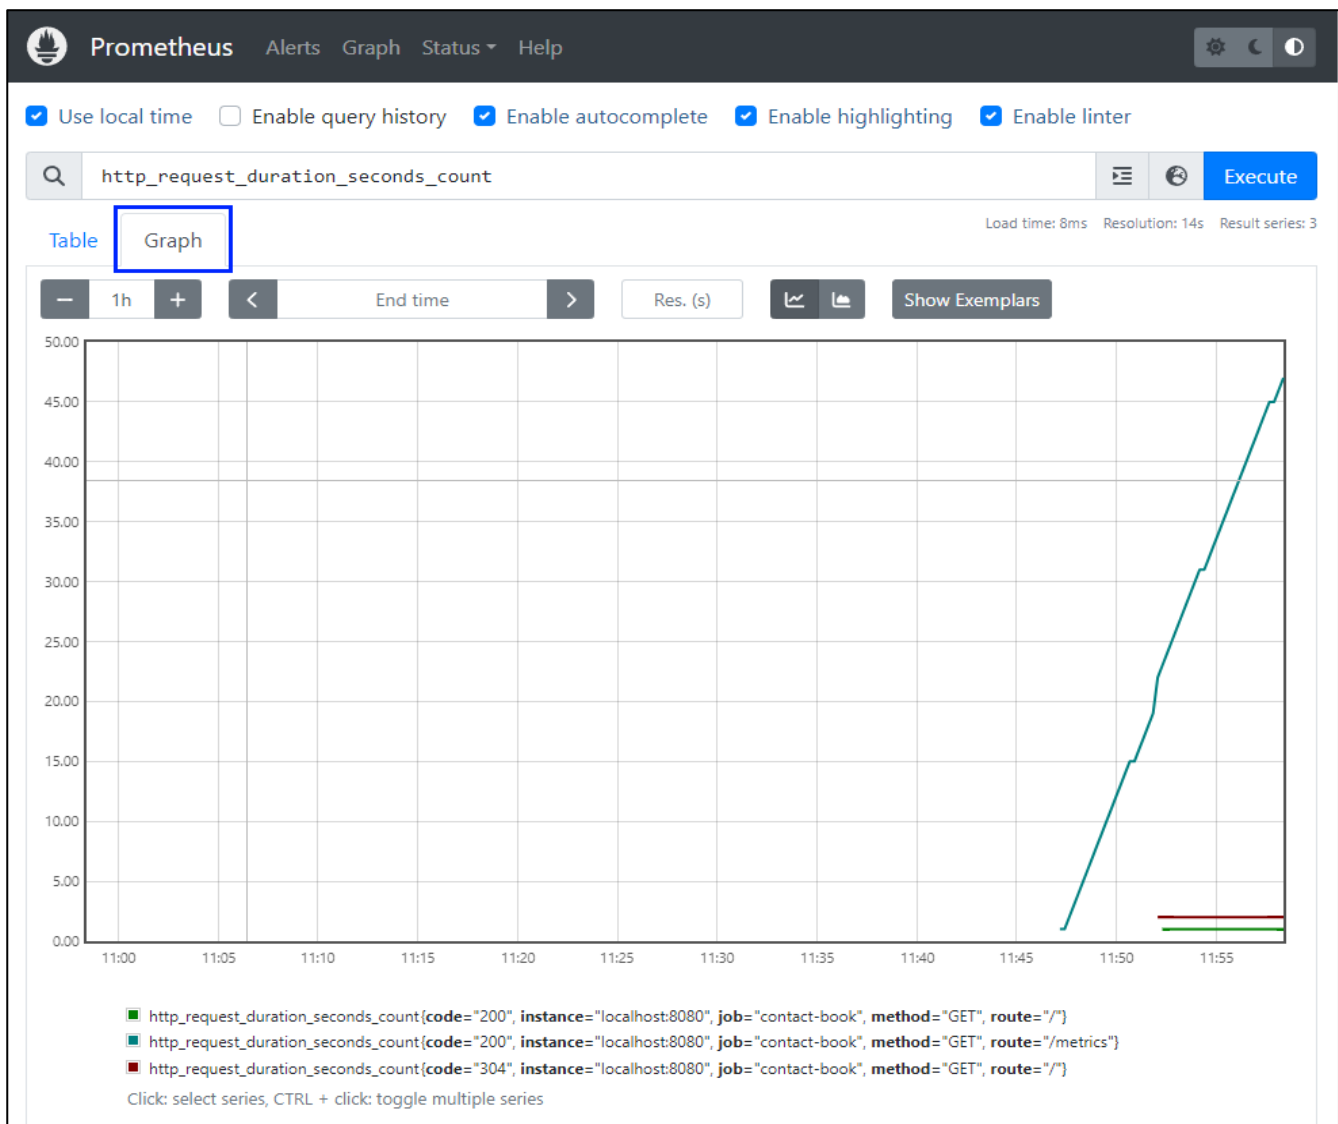
| Endpoint | State | Labels | Last Scrape | Scrape Duration | Error |
|---|-------|---|----------------|-----------------|-------|
| http://localhost:8080/metrics | UP | instance="localhost:8080" job="contact-book" | 300.000 ms ago | 16.377ms | |

NOTE: the "Contact Book" app should be running to expose metrics.

Now you can go back to the **[Graph]** page and **display some of the metric values**, using an **expression**. For example, let's see the **count of all different HTTP requests**:



Switch to **[Graph]** to look at a graph for the metric:



You can **examine more metrics** you want. When **metrics change**, click on **[Execute]** to load the changed graph.

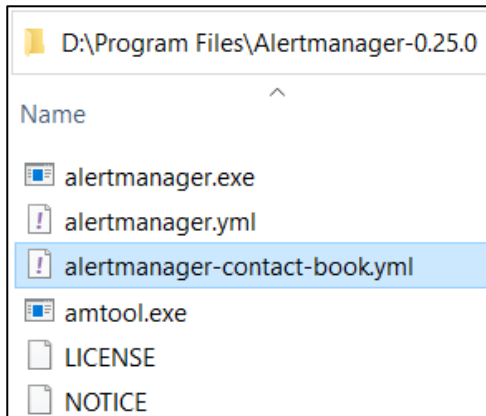
As we know how to **work with Prometheus**, let's see how to **add Alertmanager to manage alerts and send notifications**.

6. Manage "Contact Book" App Alerts with Alertmanager

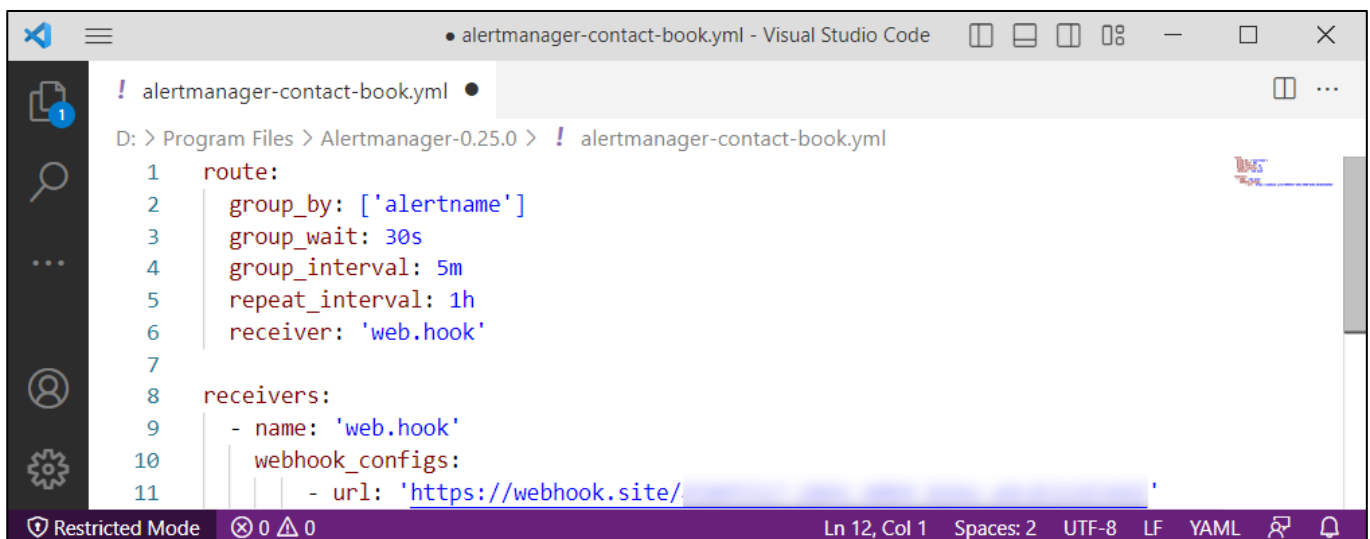
In this task, we will manage Prometheus alerts with Alertmanager and send them to Webhook.site to keep them. Our aim is to fire alerts when any page has been accessed more than 3 times during a 5-minute period.

Step 1: Configure and Run Alertmanager

Let's first see how to write a configuration file for Alertmanager to handle alerts. First, create a YAML file in the Alertmanager installation directory:



Open the file in an editor and write the sample configuration below:

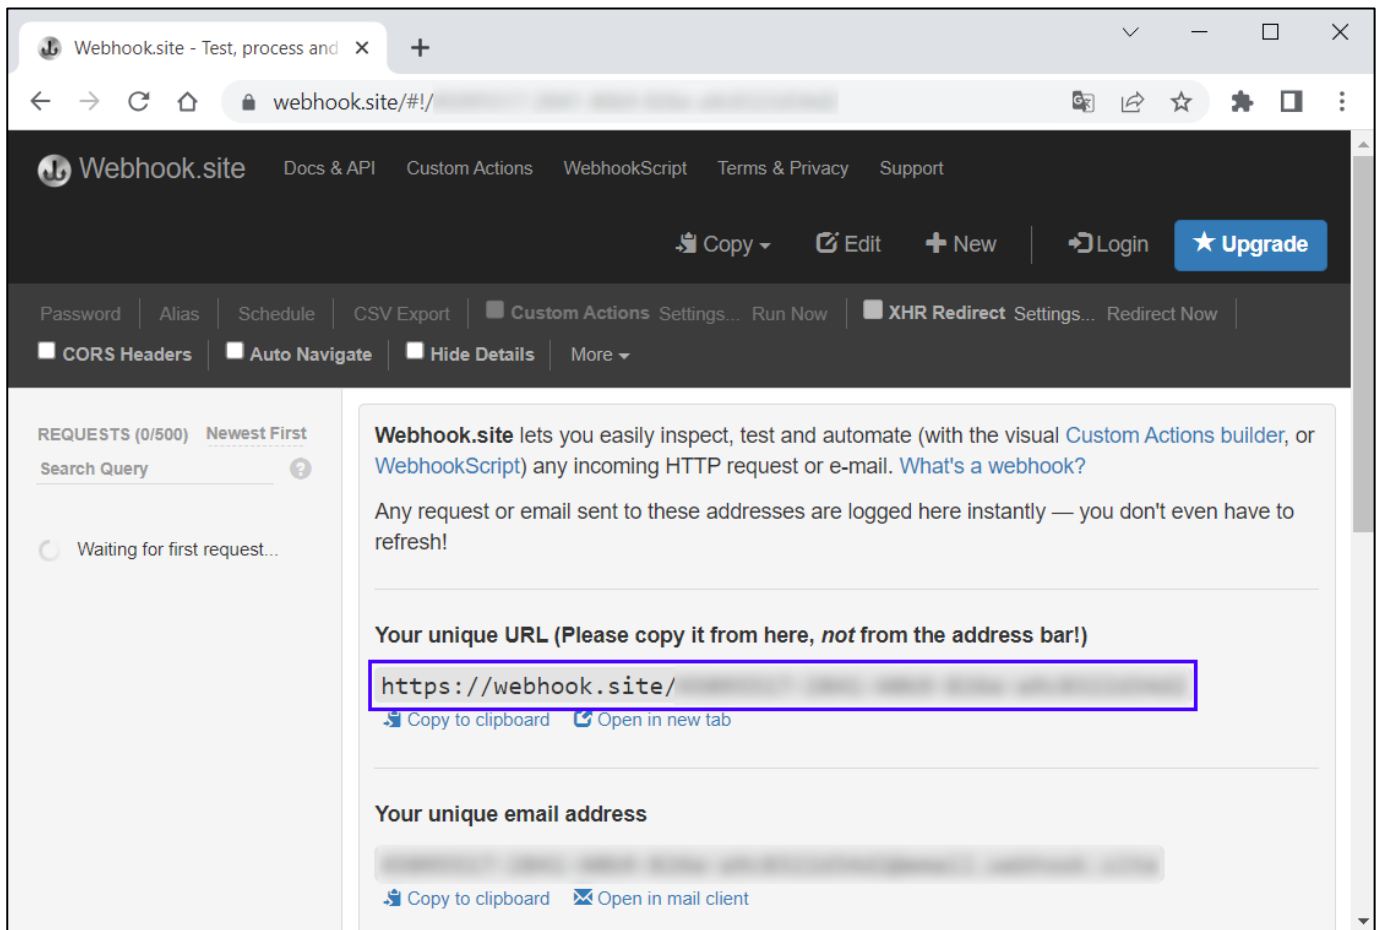


This configuration:

- Groups alerts by their name
- Sets alerts to be grouped together for a period of 30 seconds before being sent
- Sets notifications for unresolved alert groups to be sent every 5 minutes
- Sets notifications for unresolved alerts to be repeated every 1 hour
- Defines notification receiver to be "web.hook"
- Configures receiver

As you can see, the configuration contains a Webhook.site URL, which is unique. Webhook.site allows us to create temporary endpoints (webhooks) and capture the incoming requests sent to those endpoints, e.g., our Prometheus notifications.

To get your URL, navigate to Webhook.site and copy the provided URL, without closing the browser tab after this:

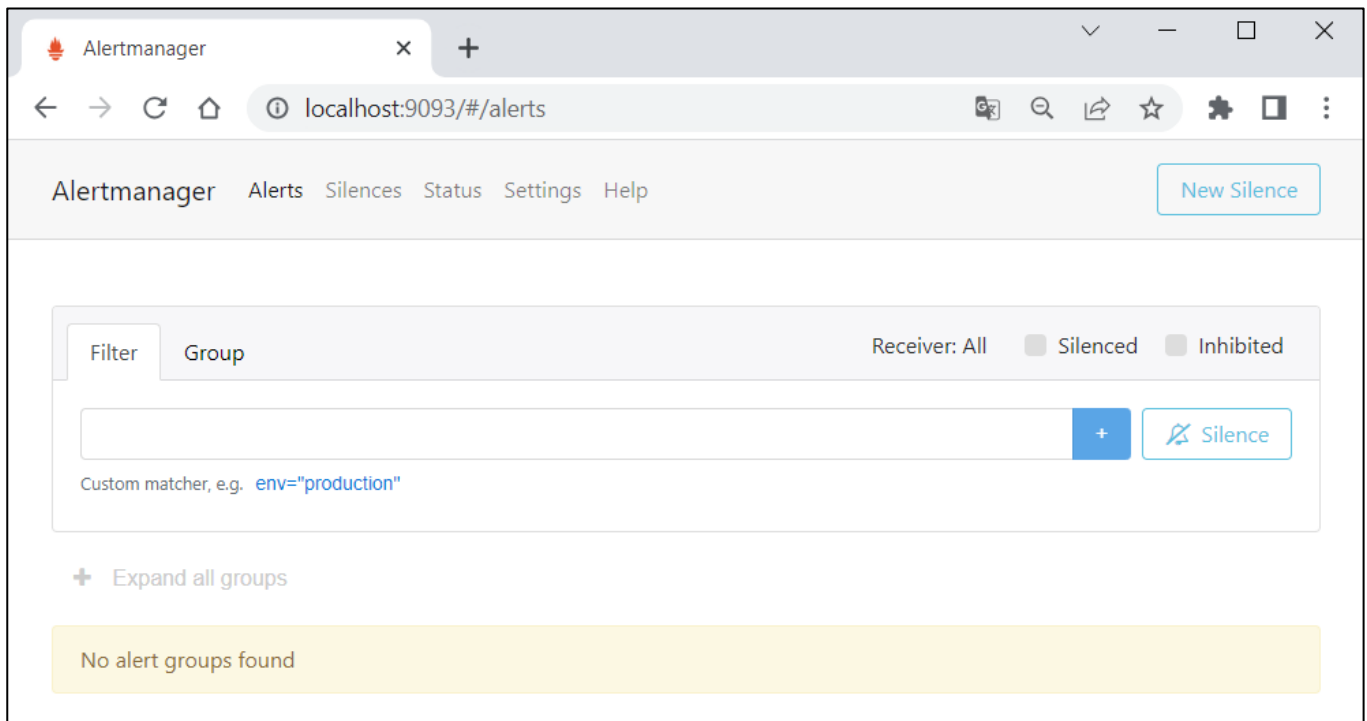


Add the URL to your configuration file and save it.

Next, run Alertmanager with the configuration file:

```
Windows PowerShell
PS D:\Program Files\Alertmanager-0.25.0> .\alertmanager --config.file .\alertmanager-contact-book.yml
ts=2023-06-02T11:23:36.408Z caller=main.go:240 level=info msg="Starting Alertmanager" version="(version=0.25.0, branch=HEAD, revision=258fab7cdd551f2cf251ed0348f0ad7289aee789)"
ts=2023-06-02T11:23:36.408Z caller=main.go:241 level=info build_context="(go=go1.19.4, user=root@0dd4f853dfffb, date=20221222-14:50:08)"
ts=2023-06-02T11:23:37.114Z caller=cluster.go:185 level=info component=cluster msg="setting advertise address explicitly" addr=fdfd::1acf:645e port=9094
ts=2023-06-02T11:23:37.121Z caller=cluster.go:681 level=info component=cluster msg="waiting for gossip to settle..." interval=2s
ts=2023-06-02T11:23:37.146Z caller=coordinator.go:113 level=info component=configuration msg="Loading configuration file" file=.\alertmanager-contact-book.yml
ts=2023-06-02T11:23:37.150Z caller=coordinator.go:126 level=info component=configuration msg="Completed loading of configuration file" file=.\alertmanager-contact-book.yml
ts=2023-06-02T11:23:37.154Z caller=tls_config.go:232 level=info msg="Listening on" address=[::]:9093
ts=2023-06-02T11:23:37.154Z caller=tls_config.go:235 level=info msg="TLS is disabled." http2=false address=[::]:9093
ts=2023-06-02T11:23:39.126Z caller=cluster.go:706 level=info component=cluster msg="gossip not settled" polls=0 before=0 now=1 elapsed=2.005274s
ts=2023-06-02T11:23:47.166Z caller=cluster.go:698 level=info component=cluster msg="gossip settled; proceeding" elapsed=10.0444373s
```

Go to <http://localhost:9093> in the browser and you should see that Alertmanager is up and working:

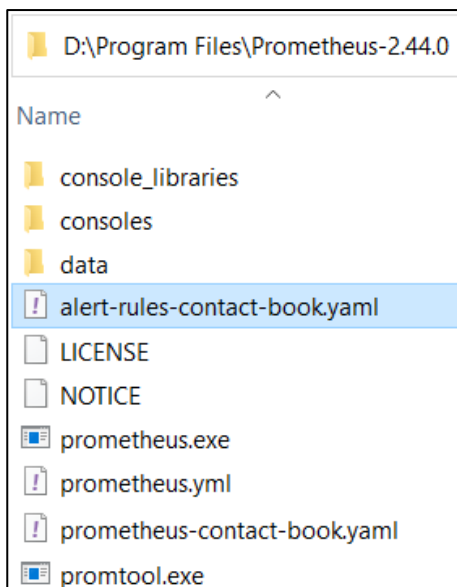


Now let's configure Prometheus to work with Alertmanager.

Step 2: Configure and Run Prometheus

We should do 2 things to make Prometheus send alerts to Alertmanager – first, create a YAML file with rules for firing an alert and, second, modify the Prometheus configuration file to use the rules and send alerts to Alertmanager.

Create a new YAML file in the Prometheus installation directory, which will define alert rules:

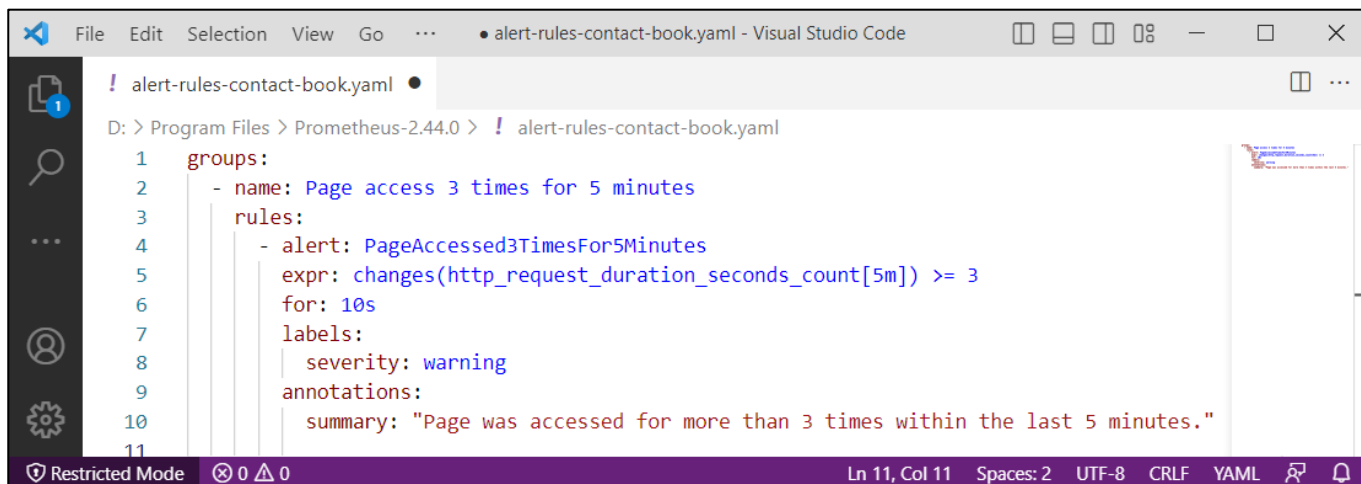


As we said earlier, we will fire an alert when an endpoint is accessed more than 3 times for 5 minutes. We will measure that using the `http_request_duration_seconds_count` metric – if its value has changed more than 3 times for the last 5 minutes. We shall have the following expression:

`changes(http_request_duration_seconds_count[5m]) >= 3`

Note that in our case we count how many times the requests count has changed on data scrape (on every 15 seconds), not how many times the count has changed generally.

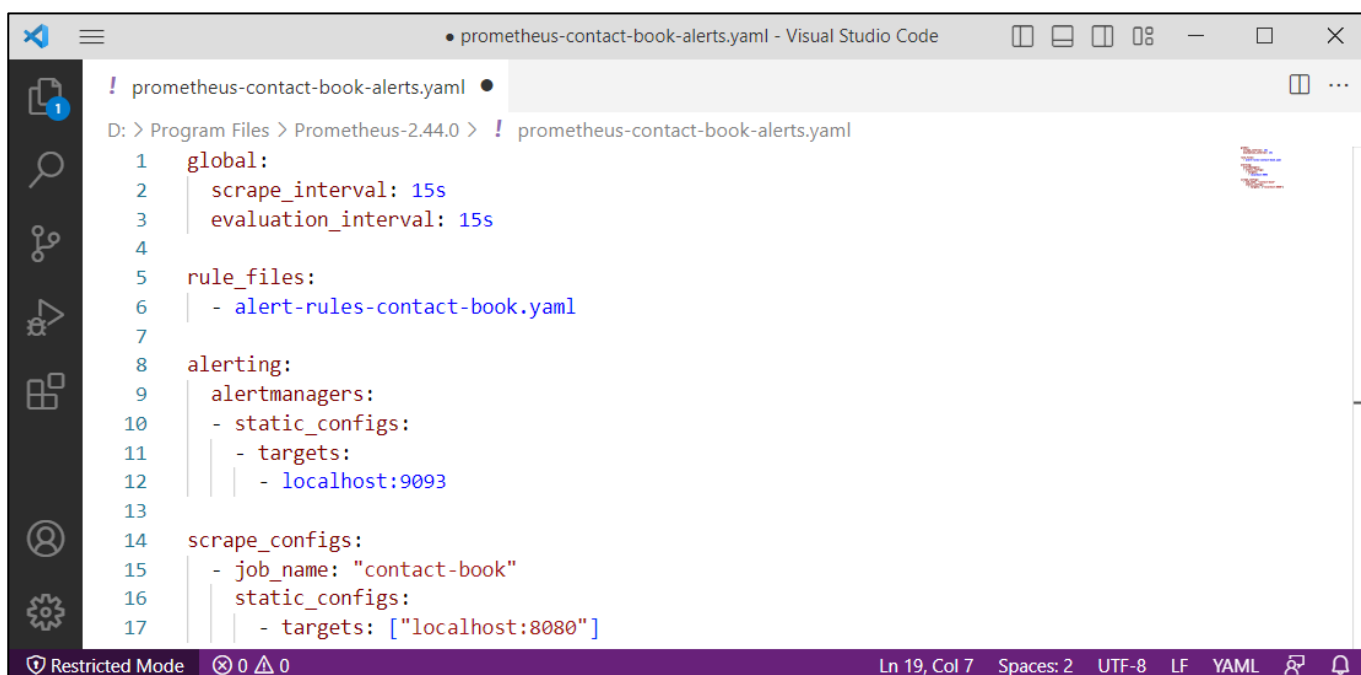

Having this expression, add the following rules configuration to the created YAML file:



```
1 groups:
2   - name: Page access 3 times for 5 minutes
3     rules:
4       - alert: PageAccessed3TimesFor5Minutes
5         expr: changes(http_request_duration_seconds_count[5m]) >= 3
6         for: 10s
7         labels:
8           severity: warning
9         annotations:
10          summary: "Page was accessed for more than 3 times within the last 5 minutes."
```

Here we have a single rules group and an alert that will be fired if the given expression is true for at least 10 seconds. The alert will have a label and summary.

Save the file and let's modify (or create a new separate file) the Prometheus configuration. It should look like this:



```
1 global:
2   scrape_interval: 15s
3   evaluation_interval: 15s
4
5 rule_files:
6   - alert-rules-contact-book.yaml
7
8 alerting:
9   alertmanagers:
10    - static_configs:
11      - targets:
12        - localhost:9093
13
14 scrape_configs:
15   - job_name: "contact-book"
16     static_configs:
17       - targets: ["localhost:8080"]
```

As you can see, we have added the name of the rules file and configurations for connection to Alertmanager, which is accessible on <http://localhost:9093> by default. Also, we have set **evaluation_interval**, which is the interval based on which Prometheus evaluates the query for alerting.

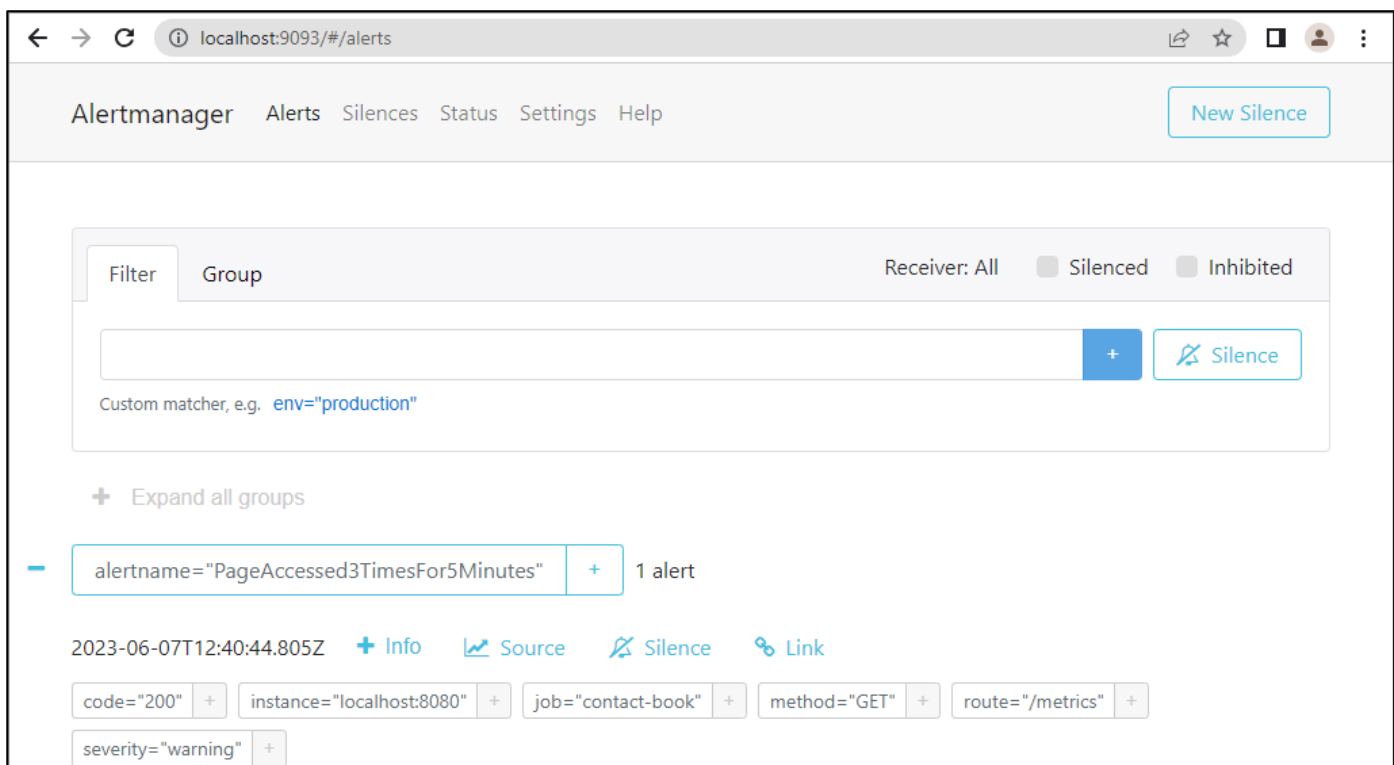
Now you should run Prometheus with the new / modified configuration file (don't forget to change the name of the configuration file if you have named it in a different way):

```
Windows PowerShell
Windows PowerShell
Copyright (C) Microsoft Corporation. All rights reserved.

Try the new cross-platform PowerShell https://aka.ms/pscore6

PS D:\Program Files\Prometheus-2.44.0> .\prometheus --config.file .\prometheus-contact-book-alerts.yaml
```

Go to the **Contact Book app** on <http://localhost:8080/> and reload the page more than 3 times. Now, access **Alertmanager** on <http://localhost:9093> and you should be able to see the new alert:



Now, visit the opened tab with Webhook.site. You should be able to see the detailed info about the sent incoming requests:

Webhook.site

Docs & APICustom ActionsWebhookScriptTerms & PrivacySupport

CopyEditNewLoginUpgrade

Request Details

POSThttps://webhook.site/...
Host...
Date06/07/2023 3:41:17 PM (19 minutes ago)
Size1.0 KB
ID5ad13a1f-452b-4107-bb70-7602d5c0f17b

Headers

connectionclose
content-typeapplication/json
content-length1043
user-agentAlertmanager/0.25.0
hostwebhook.site

Files

Query strings

Form values

Raw Content

```
{
  "receiver": "web\\.hook",
  "status": "firing",
  "alerts": [
    {
      "status": "firing",
      "labels": {
        "alertname": "PageAccessed3TimesFor5Minutes",
        "code": "200",
        "instance": "localhost:8080",
        "job": "contact-book",
        "method": "GET",
        "route": "/metrics",
        "severity": "warning"
      },
      "annotations": {
        "summary": "Page was accessed for more than 3 times within the last 5 minutes"
      },
      "startsAt": "2023-06-07T12:40:44.805Z",
      "endsAt": "0001-01-01T00:00:00Z",
      "generatorURL": "http://localhost:8080/metrics",
      "fingerprint": "0395b61aaa446a25"
    }
  ],
  "groupLabels": {
    "alertname": "PageAccessed3TimesFor5Minutes"
  },
  "commonLabels": {
    "alertname": "PageAccessed3TimesFor5Minutes",
    "code": "200",
    "instance": "localhost:8080"
  }
}
```

REQUESTS (1000) Newest First

Search Query

POST #5ad13 77.238.86.190
06/07/2023 3:41:17 PM

First -- Prev Next -- Last